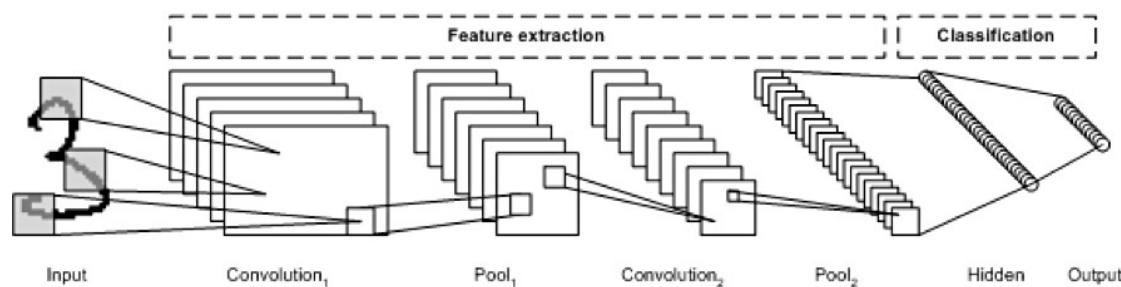


Convolutional Neural Networks



Class 7: Petia Radeva

- ↗ AI, Machine learning & Deep learning
- ↗ What is a Convolutional Neural Network?
 - ↗ Layers
 - ↗ Optimization
- ↗ Applications

Artificial Intelligence

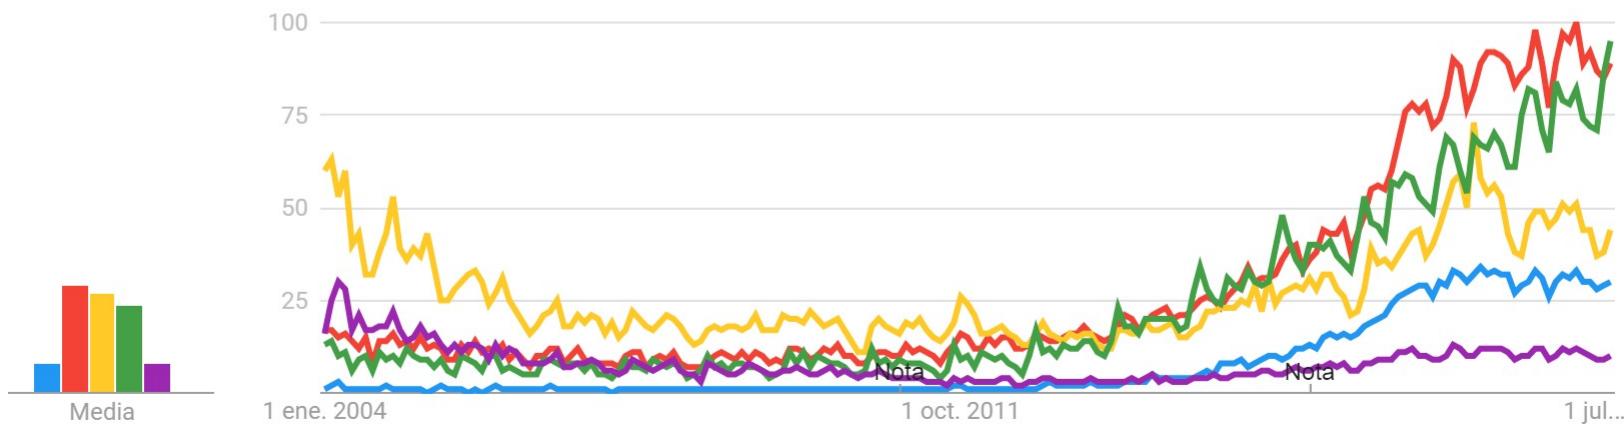
↗ Artificial Intelligence

↗ Machine Learning

↗ Deep learning

↗ Neural networks

↗ Data science



Google Scholar reveals its most influential papers

1. **"Deep Residual Learning for Image Recognition"** (2016) *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 25,256 citations
2. **"Deep learning"** (2015) *Nature* 16,750 citations
3. **"Going Deeper with Convolutions"** (2015) *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* 14,424 citations
4. **"Fully Convolutional Networks for Semantic Segmentation"** (2015) *Proceedings of the IEEE Conf. on Computer Vision and Pattern Recognition* 10,153 citations
5. **"Prevalence of Childhood and Adult Obesity in the United States, 2011-2012"** (2014) *JAMA* 8,057 citations
6. **"Global, regional, and national prevalence of overweight and obesity in children and adults during 1980-2013: a systematic analysis for the Global Burden of Disease Study 2013"** (2014) *Lancet* 7,371 citations
7. **"Observation of Gravitational Waves from a Binary Black Hole Merger"** (2016) *Physical Review Letters* 6,009 citations



Deep learning

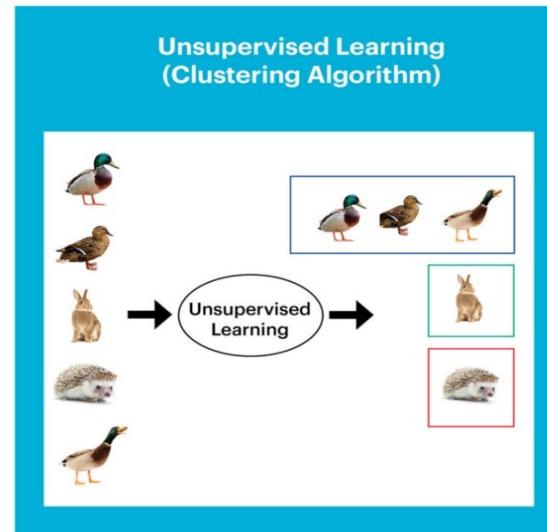
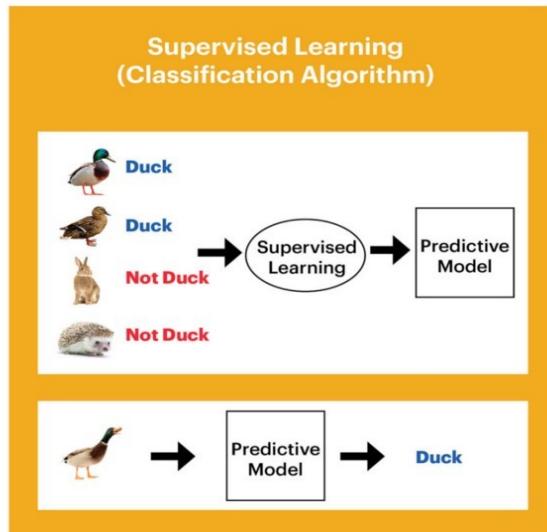
Yann LeCun, Yoshua Bengio & Geoffrey Hinton

Affiliations | Corresponding author

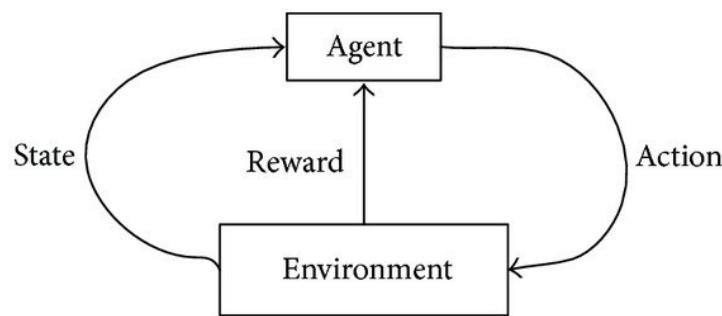
Nature 521, 436-444 (28 May 2015) | doi:10.1038/nature14539

Received: 25 February 2015 | Accepted: 01 May 2015 | Published online: 27 May 2015

Supervised vs. Unsupervised vs Reinforcement Learning

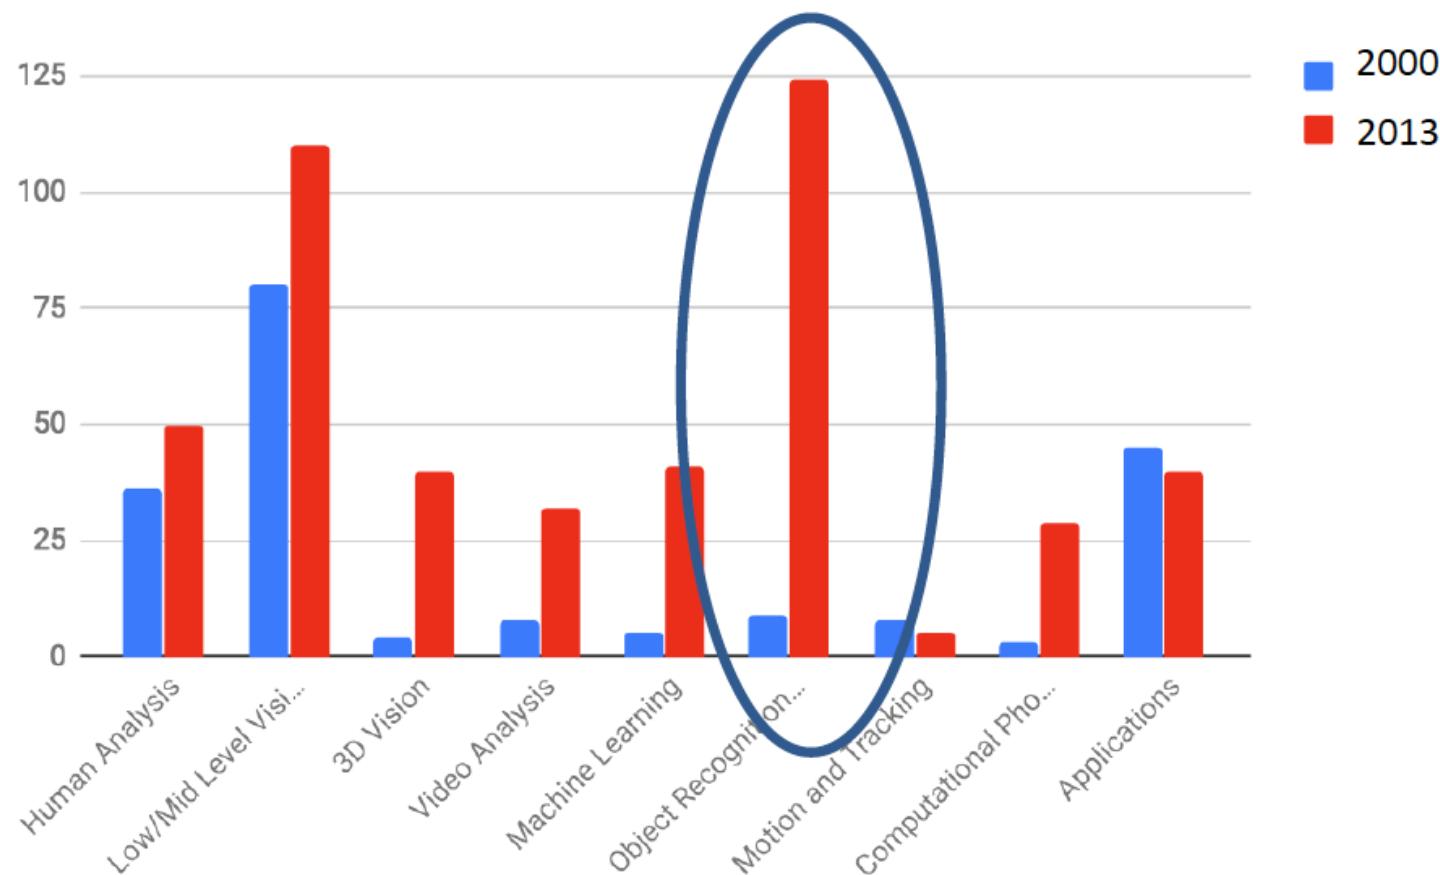


Western Digital.



CVPR topic distribution: 2000 vs 2013

CVPR topic distribution: 2000 vs. 2013



MENU ▾ Select Primary Subject Area

papers

Oral

Poster

1	Recognition: detection, categorization, retrieval	177	25	152
2	Image and video synthesis and generation	157	26	131
3	3D from multi-view and sensors	137	26	111
4	Low-level vision	110	19	91
5	Vision + language	105	20	85
6	Segmentation, grouping and shape analysis	99	16	83
7	Transfer/low-shot/long-tail learning	86	15	71
8	Deep learning architectures and techniques	85	20	65
9	Self-& semi-& meta- & unsupervised learning	84	7	77
10	Video analysis and understanding	77	15	62
11	Pose estimation and tracking	62	14	48
12	Representation learning	61	11	50
13	3D from single images	60	10	50
14	Scene analysis and understanding	56	9	47
15	Face and gestures	54	7	47
16	Computational photography	53	10	43
17	Motion and tracking	53	8	45
18	Adversarial attack and defense	52	10	42
19	Datasets and evaluation	52	7	45
20	Machine learning	41	7	34
21	Action and event recognition	40	8	32
22	Efficient learning and inferences	40	3	37
23	Medical, biological and cell microscopy	37	5	32
24	Vision applications and systems	32	4	28
25	Navigation and autonomous driving	31	2	29
26	Vision + graphics	24	6	18
27	Privacy and federated learning	21	3	18
28	Vision + X	20	4	16
29	Physics-based vision and shape-from-X	16	2	14
30	Robot vision	16	3	
31	Explainable computer vision	15	2	13
32	Demo	15		15
33	Optimization methods	14	2	12
34	Transparency, fairness, accountability, privacy and ..	14	4	
35	Document analysis and understanding	12		12
36	Biometrics	11	2	9

Categories at CVPR 2022 ranked by number of papers accepted

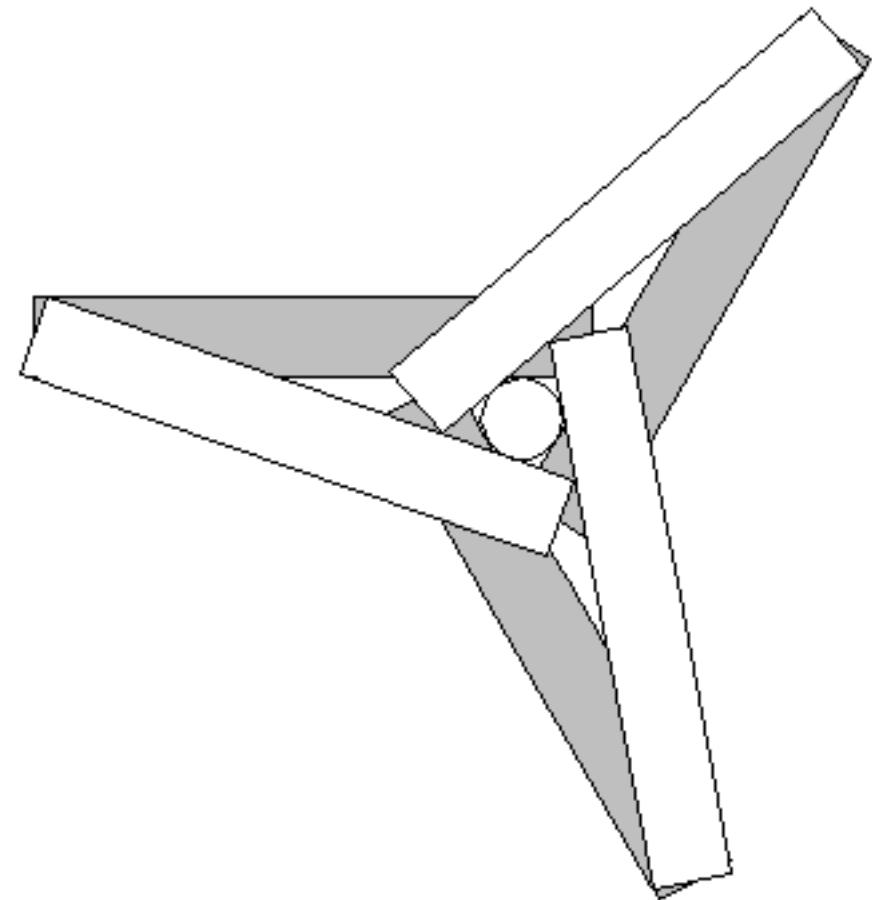
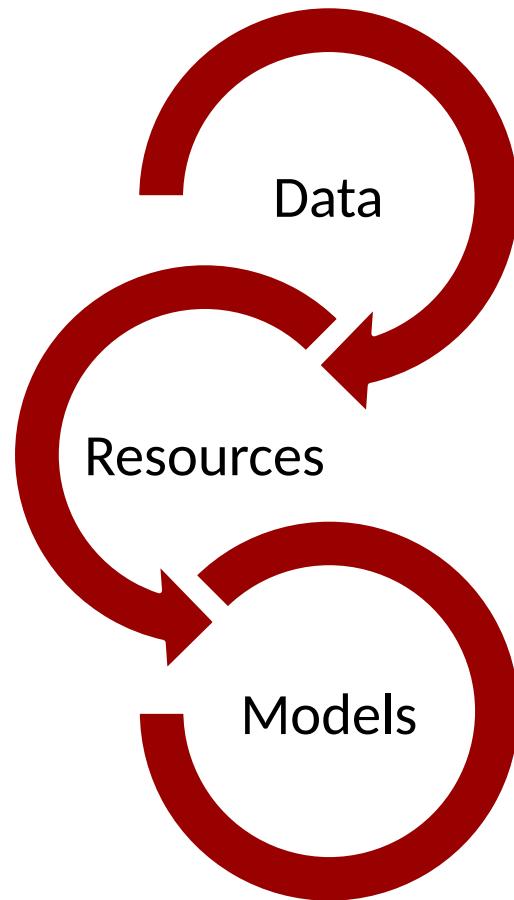


Deep learning - where does it come from?

- ↗ **1943:** neurophysiologist Warren McCulloch and mathematician Walter Pitts – a **neuronal model** using an electrical circuit
- ↗ **1950:** Alan Turing created the world-famous **Turing Test:** a computer to pass, has to be able to convince a human that it is a human and not a computer.
- ↗ **1952:** Arthur Samuel created the first computer **program which could learn** as it ran. It was a game which played checkers.
- ↗ **1958:** Frank Rosenblatt designed the first artificial neural network, called **Perceptron**. The main goal of this was pattern and shape recognition.
- ↗ **1959**, Bernard Widrow and Marcian Hoff created
 - ↗ ADELINe to detect binary patterns (in a stream of bits, to predict the next one).
 - ↗ MADELINE to eliminate echo on phone lines, still in use today.
- ↗ **1986:** Neural networks use **back propagation** allowing multiple layers to be used in a neural network, creating what are known as 'slow learners'.
- ↗ Late **1980s** and **1990s** did not bring much to the field.
- ↗ **1997**, the IBM computer Deep Blue (a chess-playing computer) beat the world chess champion.
- ↗ **1998** AT&T Bell Laboratories on digit recognition resulted in good accuracy in detecting handwritten postcodes from the US Postal Service.

What changed today?

The magic triangle



Data

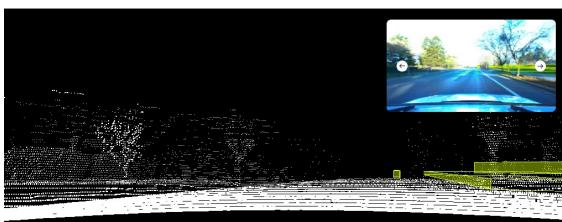
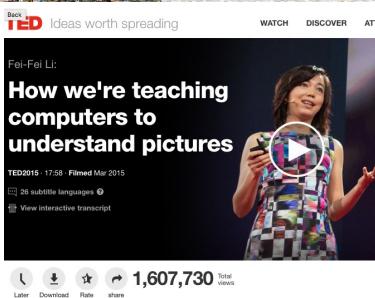
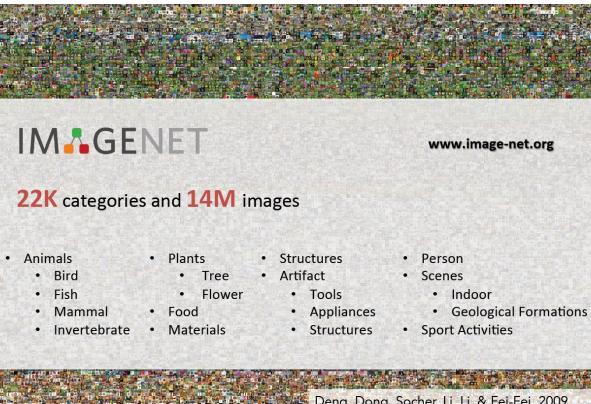
90% of all digital data were generated last 2 years.

Daily:

- ↗ 300M photos get uploaded
 - ↗ 95M photos and videos are shared on Instagram
 - ↗ 100M people use the Instagram “stories”
 - ↗ 15K GIFs are sent via Facebook
 - ↗ 154K calls on Skype
 - ↗ 4.7T photos stored in cameras



DL Datasets



Lyft Level 5

<https://www.datasetlist.com>



LVIS Challenge:
2.2M masks, 16K images

00:29 → 00:37 Steven went, got the keys and we gonna have them back. That easy.
(serious face)

00:37 → 00:40 I couldn't...
(serious face)

00:40 → 00:42 (serious face)
You said you were going to do it and you are not doing it!

Q1: How is the man who is not being blamed responding to the situation? *advanced*
A1: He thinks the other man is slackin even if he is not saying it. *intermediate*
A2: He is showing support for the woman by taking her side. *intermediate*
A3: He thinks he is better than both of the people arguing. *easy*
A4: The man is trying to ignore the woman. *intermediate*

Q2: How is the woman who is being blamed responding to the situation? *advanced*
A1: Because a small problem became a huge problem. *intermediate*
A2: She has too much on her plate, and this new problem overwhelms her. *advanced*
A3: The man is ignoring her. *intermediate*
A4: Because both of them seem to be ignoring her. *intermediate*

SocialIQ



Places2: 10M images



TACO: Waste
in the wild

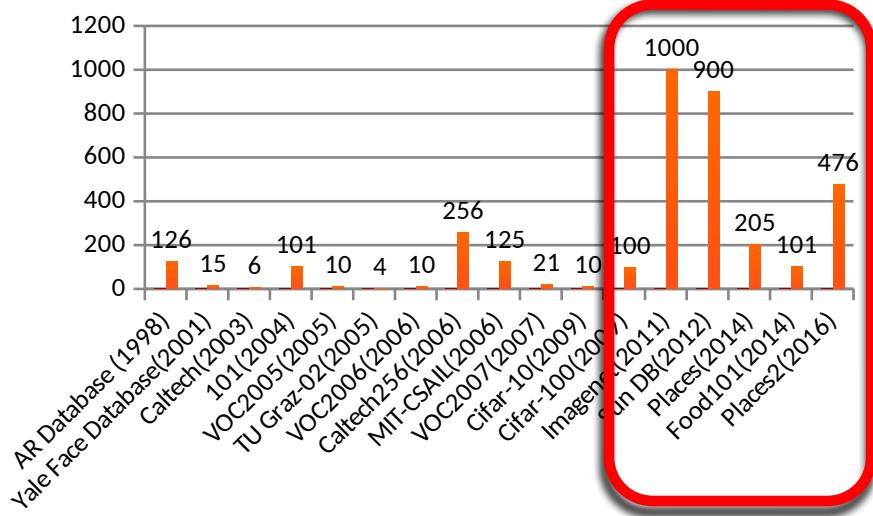


FastMRI

01:42

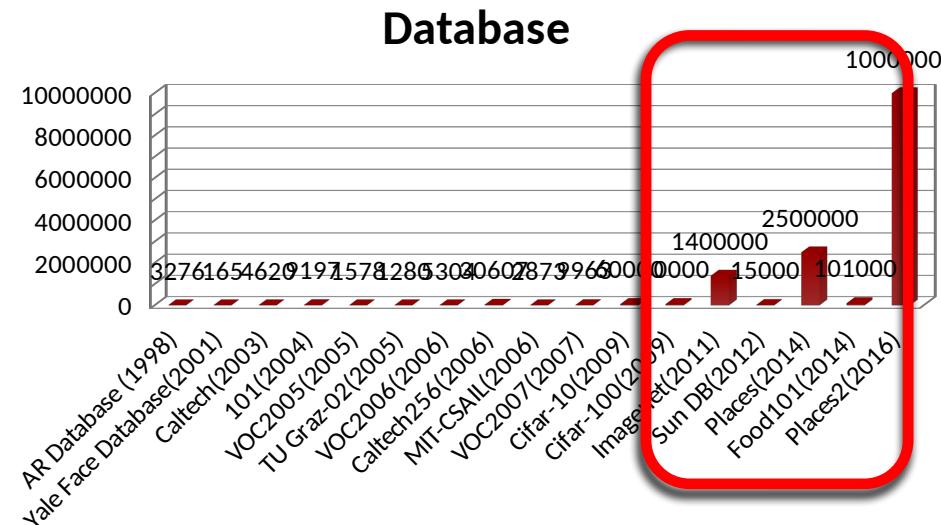
Image databases evolution

Number of objects/Database



ImageNet & Deep learning

Number of images/Database



Imagenet



IM^{GENET}

www.image-net.org

22K categories and **14M** images

- Animals
 - Bird
 - Fish
 - Mammal
 - Invertebrate
- Plants
 - Tree
 - Flower
 - Food
 - Materials
- Structures
 - Artifact
 - Tools
 - Appliances
 - Structures
- Person
- Scenes
 - Indoor
 - Geological Formations
 - Sport Activities



Deng, Dong, Socher, Li, Li, & Fei-Fei, 2009

Back TED Ideas worth spreading

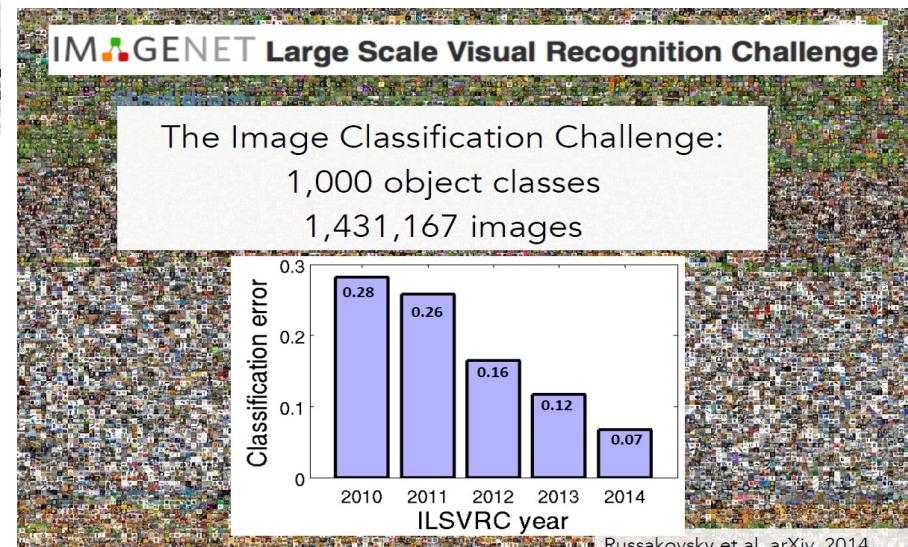
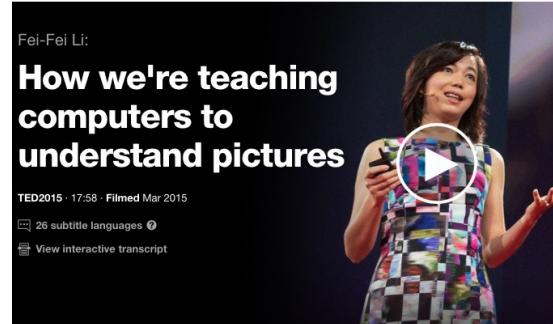
WATCH DISCOVER ATT

Fei-Fei Li:
How we're teaching computers to understand pictures

TED2015 · 17:58 · Filmed Mar 2015

26 subtitle languages ⓘ View interactive transcript

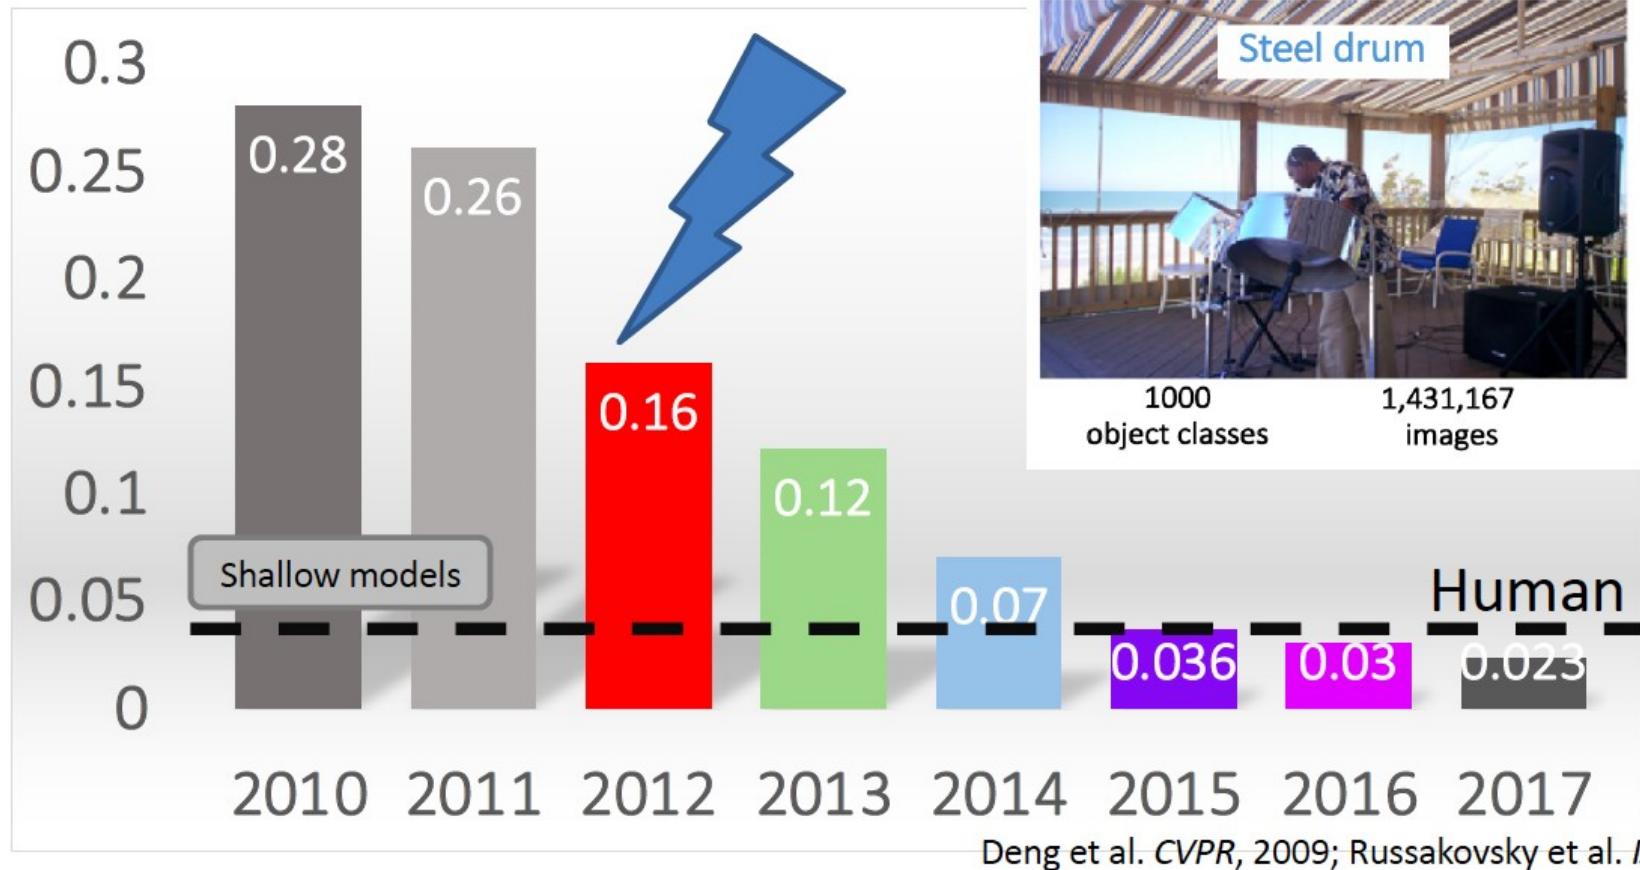
Later Download Rate share 1,607,730 Total views





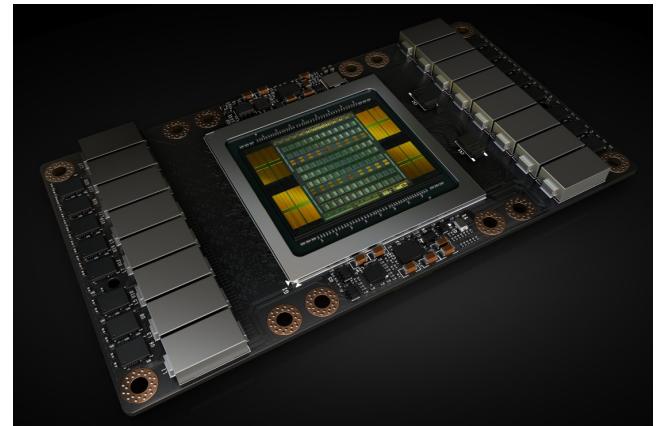
Classification Task

Classification Error



The Importance of GPUs

- ↗ Nvidia Tensor Cores - 2017
- ↗ Google Tensor Processing Unit (TPU) - 2016
- ↗ Intel - Nervana Neural Processor - 2017
- ↗ GPUs in Cloud Computing (Google, 2017)



$$D = \begin{pmatrix} A_{0,0} & A_{0,1} & A_{0,2} & A_{0,3} \\ A_{1,0} & A_{1,1} & A_{1,2} & A_{1,3} \\ A_{2,0} & A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,0} & A_{3,1} & A_{3,2} & A_{3,3} \end{pmatrix}_{\text{FP16 or FP32}} \begin{pmatrix} B_{0,0} & B_{0,1} & B_{0,2} & B_{0,3} \\ B_{1,0} & B_{1,1} & B_{1,2} & B_{1,3} \\ B_{2,0} & B_{2,1} & B_{2,2} & B_{2,3} \\ B_{3,0} & B_{3,1} & B_{3,2} & B_{3,3} \end{pmatrix}_{\text{FP16}} + \begin{pmatrix} C_{0,0} & C_{0,1} & C_{0,2} & C_{0,3} \\ C_{1,0} & C_{1,1} & C_{1,2} & C_{1,3} \\ C_{2,0} & C_{2,1} & C_{2,2} & C_{2,3} \\ C_{3,0} & C_{3,1} & C_{3,2} & C_{3,3} \end{pmatrix}_{\text{FP16 or FP32}}$$

GPU cores is based on matrix multiplication

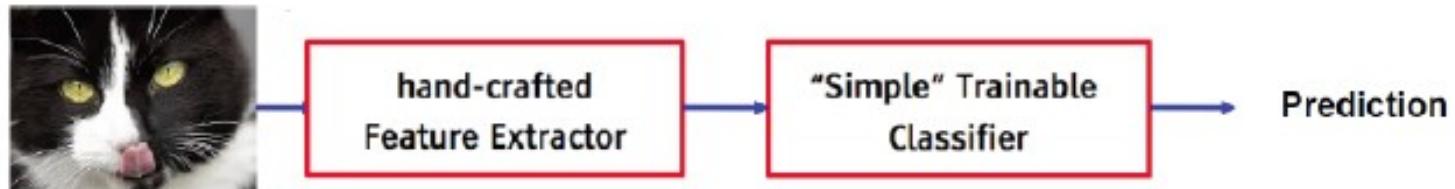


What is a Neural Network?

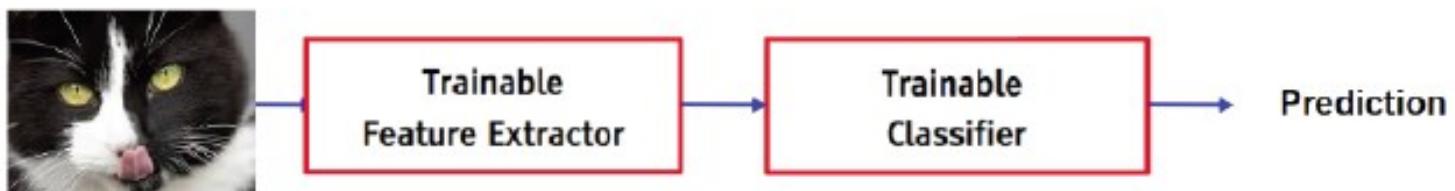
- ↗ AI, Machine learning & Deep learning
- ↗ What is a Convolutional Neural Network?
 - ↗ Layers
 - ↗ Optimization
- ↗ Applications

Why Deep learning?

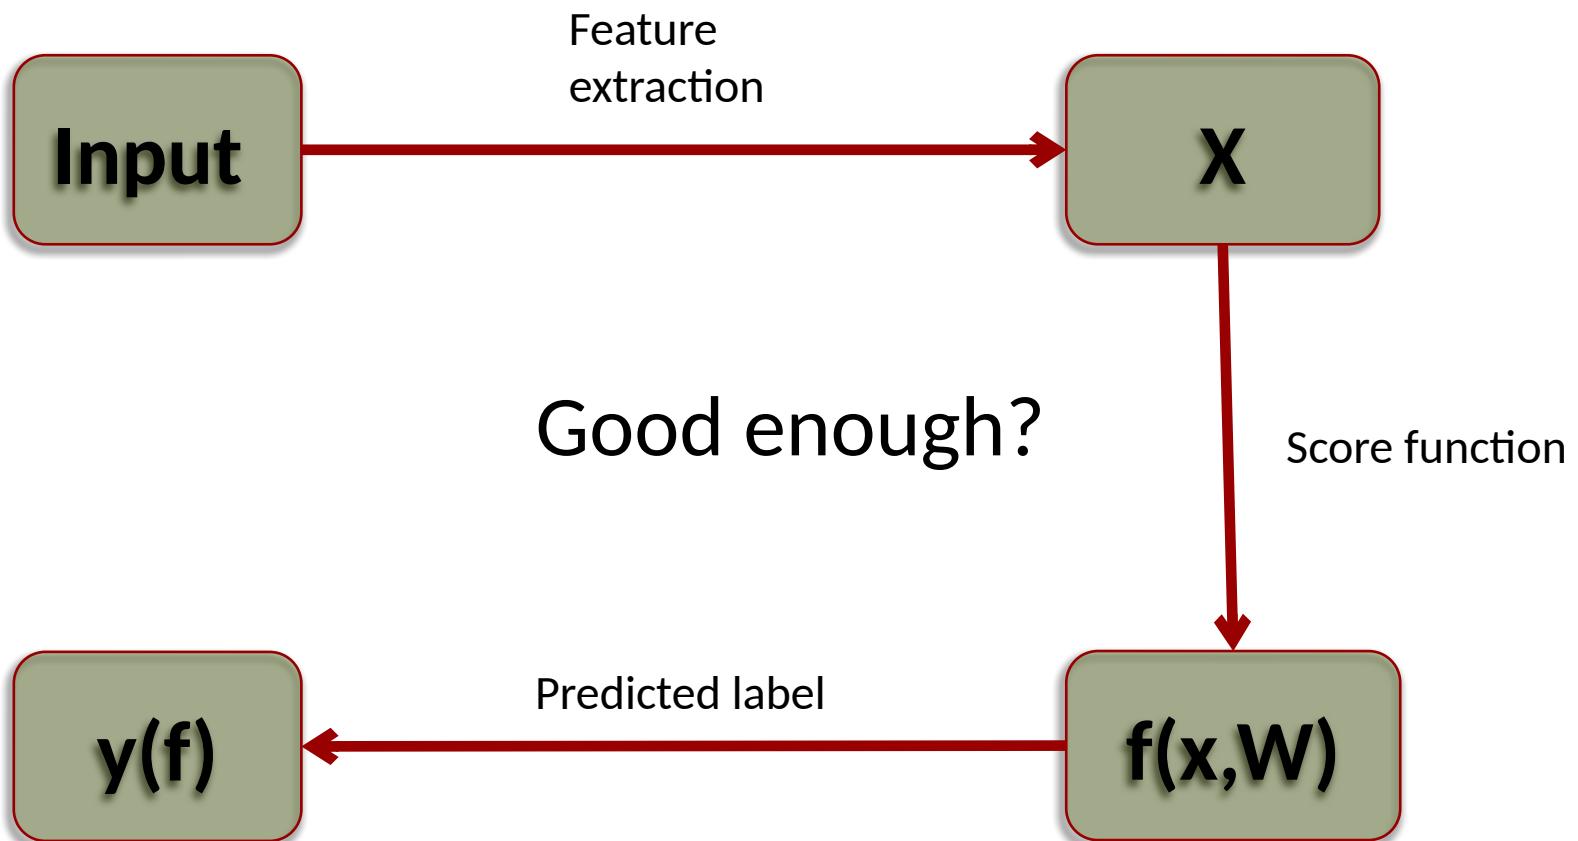
Classical way of solving Computer Vision problems



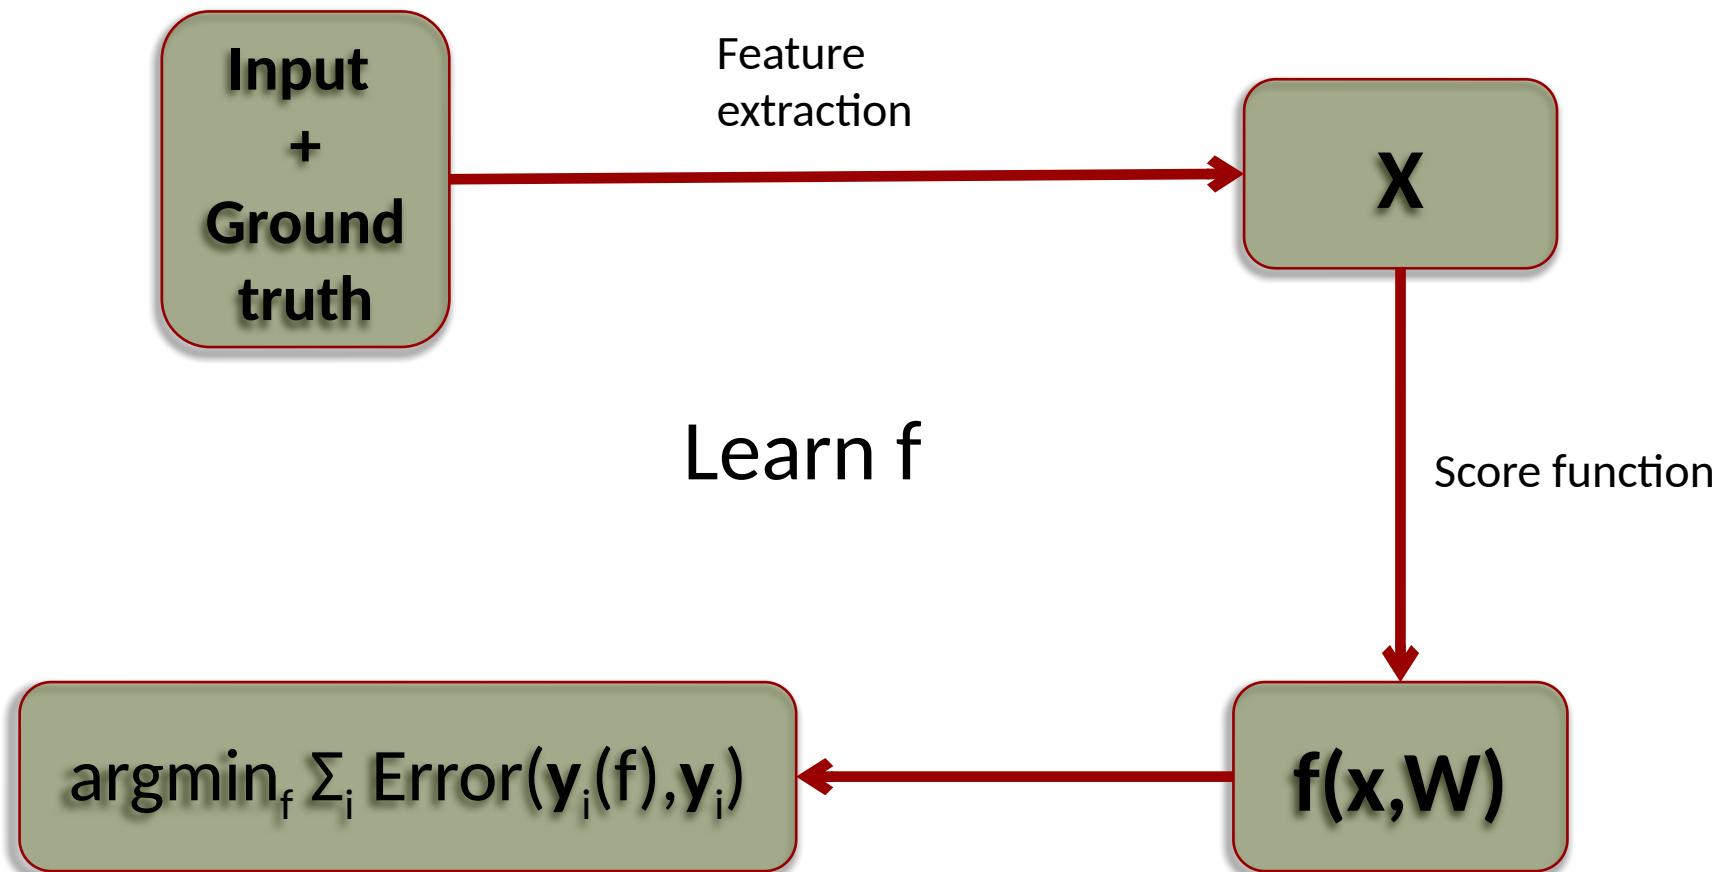
How Computer Vision problems are solved by Deep Learning



The learning pipeline



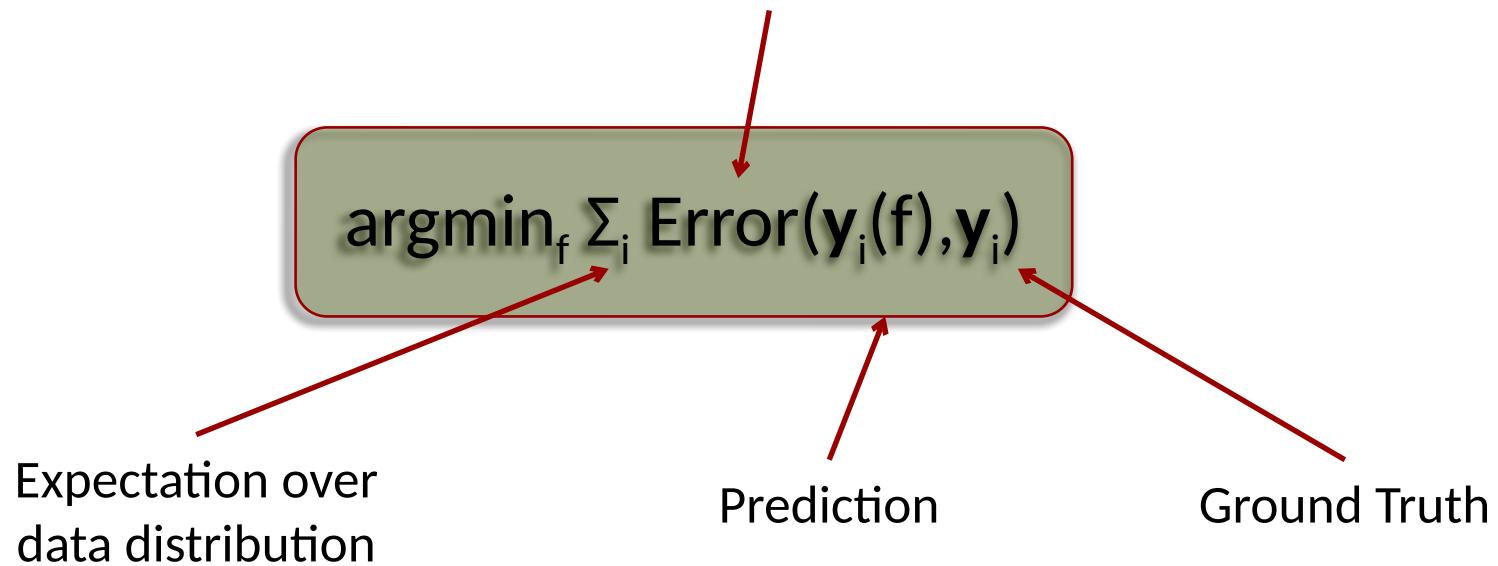
The training process



The learning process

Training data $\{(x_i, y_i), i = 1, 2, \dots, n\}$

Measure of prediction quality (error, loss)

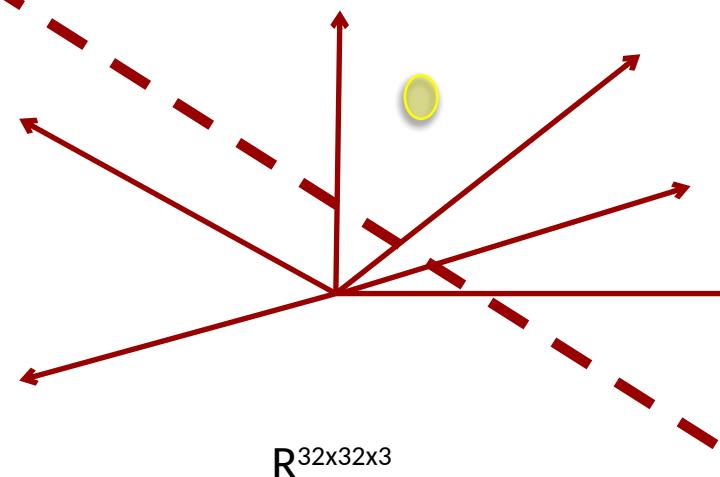


Loss function the negative conditional log-likelihood, with the interpretation that $f_i(X)$ estimates $P(Y=i | X)$:

$$L(f(x), y) = -\log f_i(x), \text{ where } f_i(x) \geq 0, \sum_i f_i(x) = 1.$$

Linear classification

Given two classes how to learn a hyperplane to separate them?



To find the hyperplane that separates dogs from cats, we need to define:

- The score function
- The loss function
- And the optimization process.

How to project data in the feature space:

3x1

3072x1

$$f(x) = W x + b$$

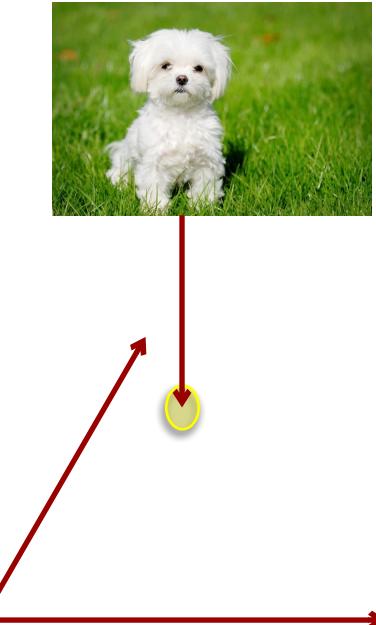
3x3072

3x1

If x is an image of $(32 \times 32 \times 3)$, $\rightarrow x$ in R^{3072} ,

The matrix W is (3×3072) .

The bias vector b is 3-dimensional.



How to project data in the feature space:

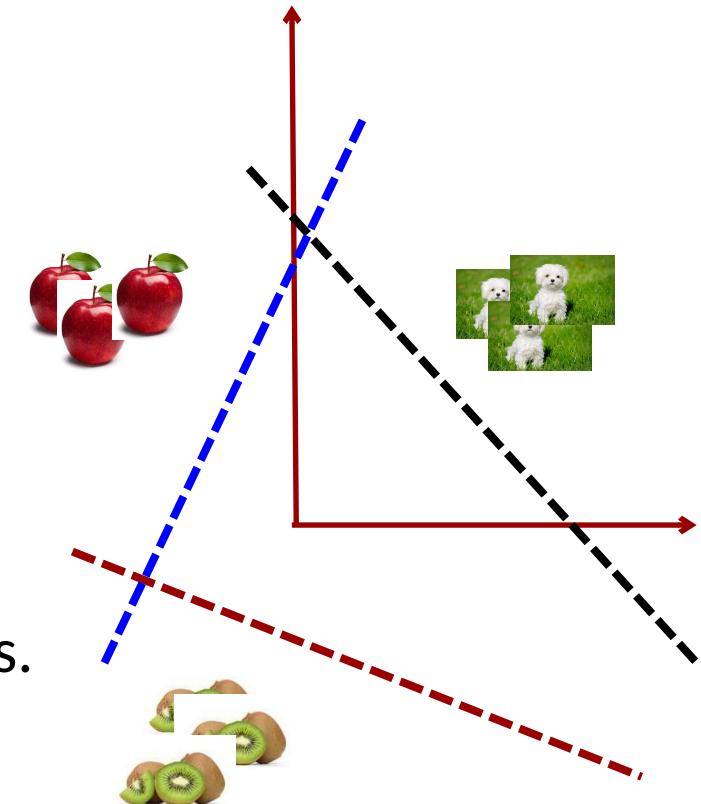
3x1

3072x1

$$f(x) = W x + b$$

3x3072

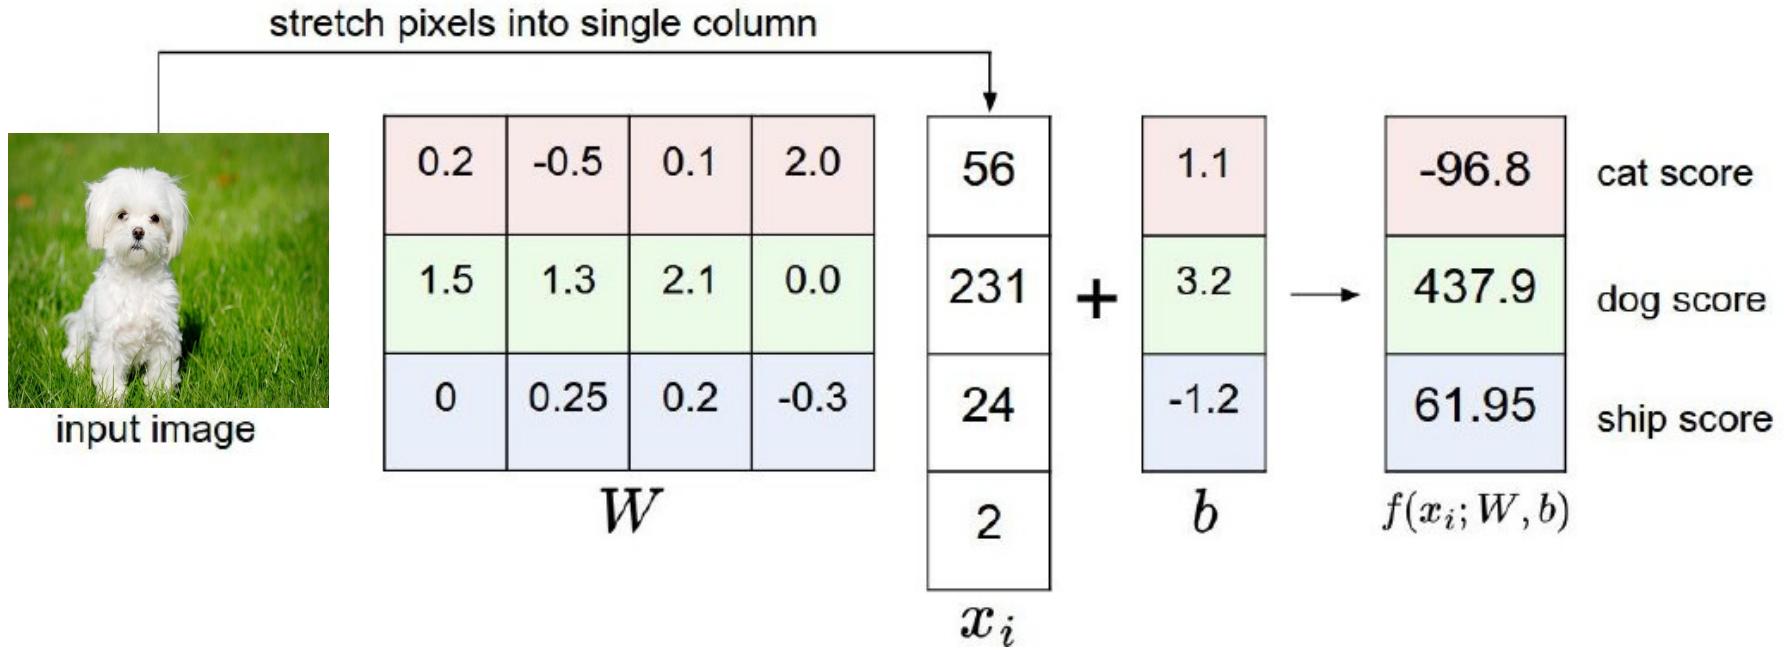
3x1



If we have 3 classes, $f(x)$ will give 3 scores.

Image classification

Example with an image with 4 pixels, and 3 classes (**cat/dog/ship**)

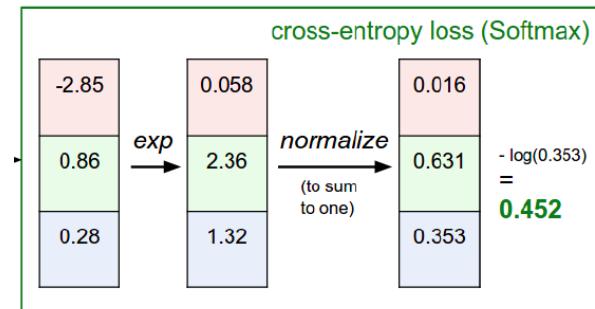


Loss function and optimisation

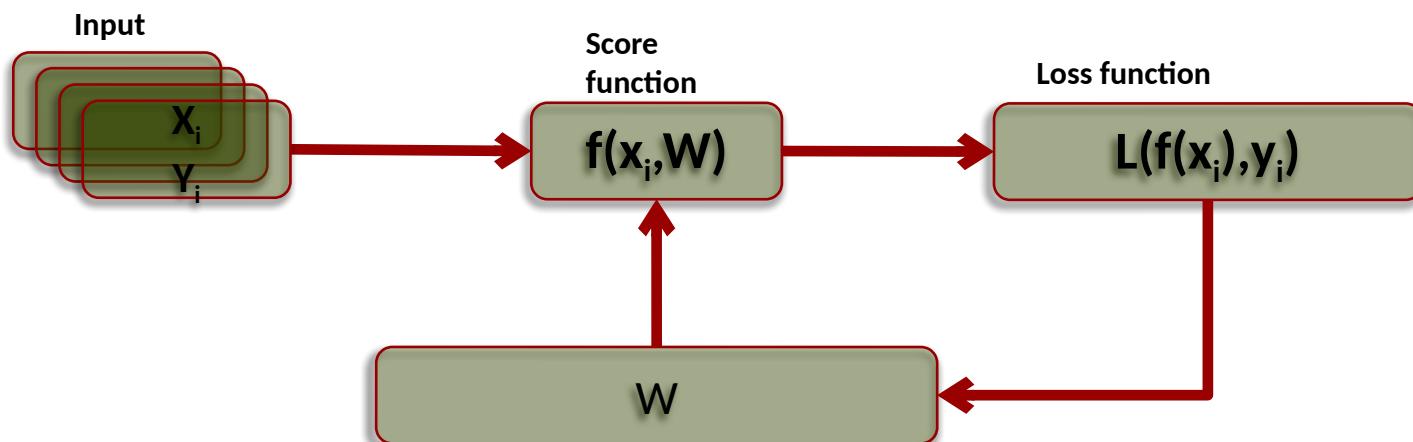
↗ **Question:** if you were to assign a single number to how unhappy you are with these scores, what would you do?

$$L_i = -\log \left(\frac{e^{f_{y_i}}}{\sum_j e^{f_j}} \right)$$

softmax function



Question : Given the score and the loss function, how to find the parameters W?



- ↗ AI, Machine learning & Deep learning
- ↗ What is a Convolutional Neural Network?
 - ↗ Layers
 - ↗ Optimization
- ↗ Applications

How is a CNN doing deep learning?

$$y = Wx$$



$$\begin{matrix} 0.2 & -0.5 & 0.1 & 2.0 \\ 1.5 & 1.3 & 2.1 & 0.0 \\ 0 & 0.25 & 0.2 & -0.3 \end{matrix}$$

W

$$y_1 = \sum_i W_{1i} x_i$$

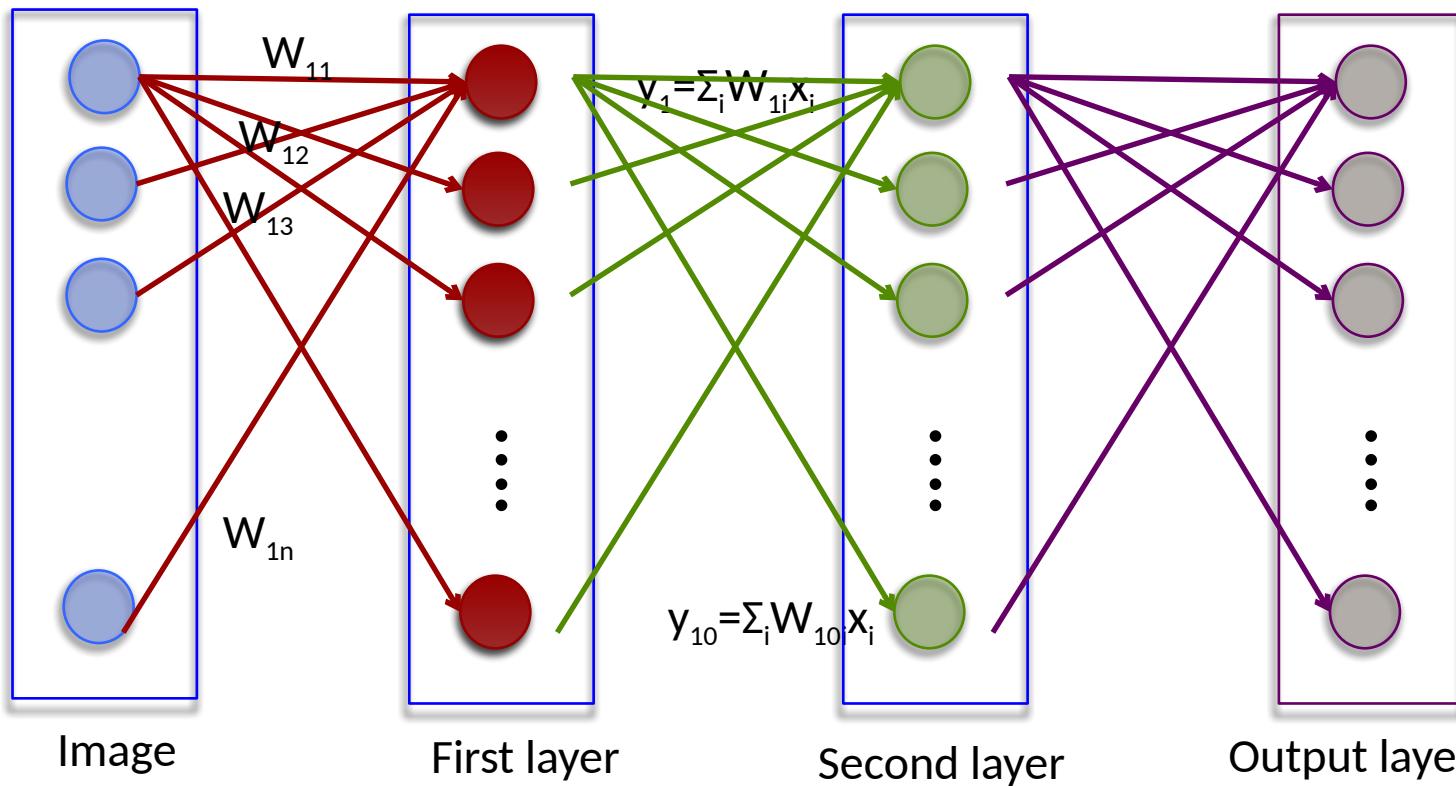
$$\begin{array}{c} 56 \\ 231 \\ 24 \\ 2 \end{array} + \begin{array}{c} 1.1 \\ 3.2 \\ -1.2 \end{array} \rightarrow \begin{array}{c} \dots \\ -96.8 \\ 437.9 \\ 61.95 \end{array}$$

x_i b $f(x_i; W, b)$

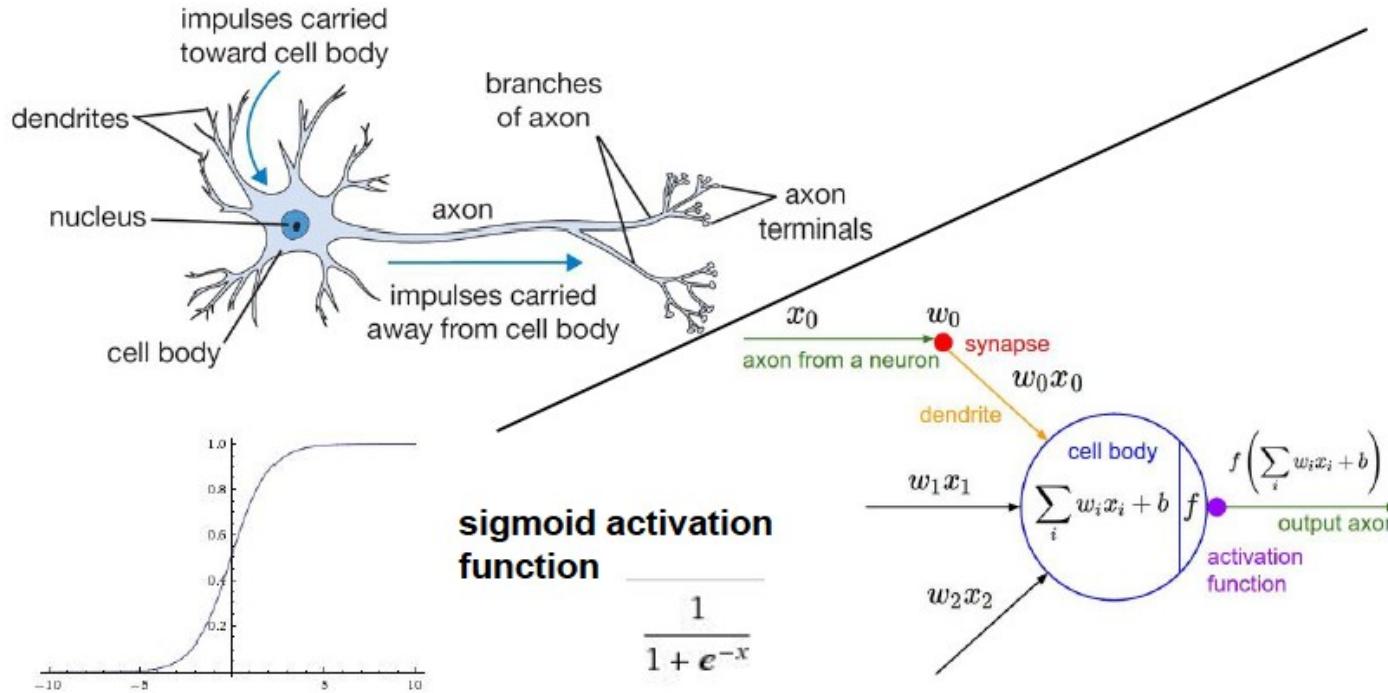
$$y = W(Wx)$$

$$y = W(W(Wx))$$

Fully connected layers



Why a CNN is a neural network?

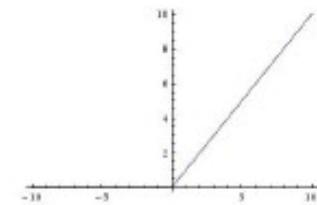


Modern CNNs – 10M neurons

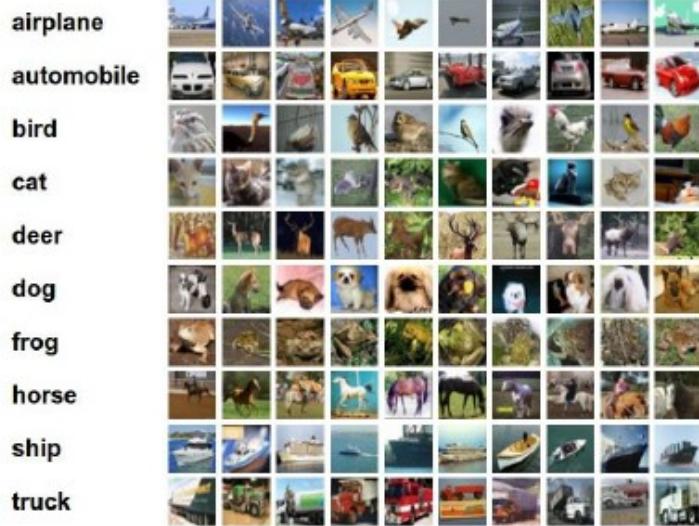
Human CNNs – 5B of neurons.

From: Fei-Fei Li & Andrej Karpathy & Justin Johnson

ReLU $\max(0, x)$



Why is it convolutional?



$$f(x_i; W, b) = Wx_i + b$$

Diagram illustrating the computation of a feature map element:

The input image x_i is represented by a 4x4 grid of values. The weight matrix W is a 4x4 matrix. The bias vector b is a 4-element vector. The output value is calculated as the dot product of the row of W and the column of x_i , plus the corresponding element of b .

0.2	-0.5	0.1	2.0	56	1.1	-96.8
1.5	1.3	2.1	0.0	231	3.2	437.9
0	0.25	0.2	-0.3	24	-1.2	61.95
				2		

W

x_i

b

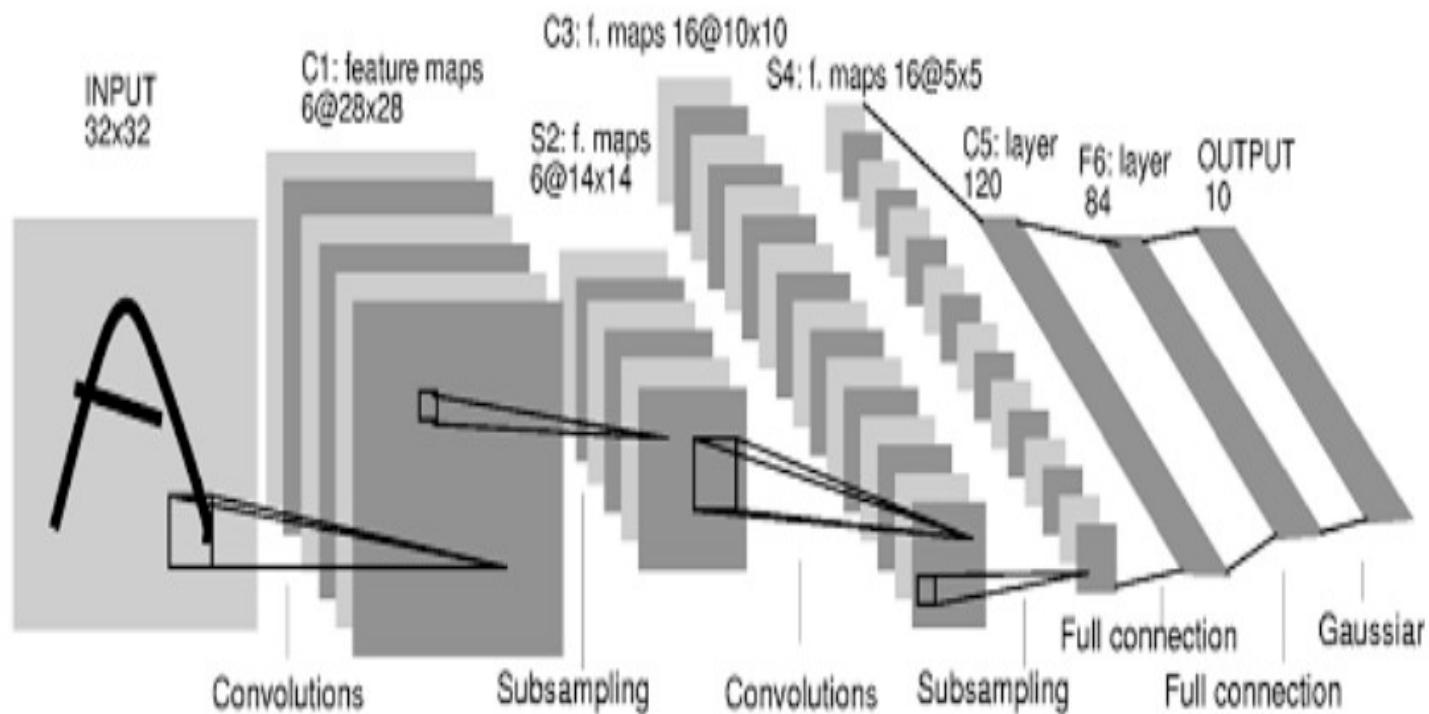
$f(x_i; W, b)$

Adapted from: Fei-Fei Li & Andrej Karpathy & Justin Johnson

What is new in the Convolutional Neural Network?

1998

LeCun et al.



CNN evolution

lecture1.pdf - Adobe Acrobat Reader DC

Archivo Edición Ver Ventana Ayuda

Inicio Herramientas lecture1.pdf x

33 / 45 | 111% |

1998 LeCun et al.

The diagram illustrates a convolutional neural network architecture. It starts with an input image of size 32x32. This is processed through two convolutional layers (C1 and C2) with feature maps of sizes 6x6 and 16x16 respectively, followed by two subsampling layers (S2 and S4). The output is then processed through two more convolutional layers (C3 and C4) with feature maps of sizes 16x10 and 16x5 respectively, followed by two subsampling layers (S3 and S5). The final output is a single value produced by a Gaussian layer. The diagram also shows a handwritten digit 'A' as input.

of transistors 10^6

of pixels used in training 10^7

2012 Krizhevsky et al.

The diagram illustrates the AlexNet architecture, a deep learning model. It takes an input image of size 224x224 and processes it through two sets of convolutional and pooling layers. The first set consists of five layers: two convolutional layers with 11x11 kernels and stride of 4, followed by two max pooling layers with 55x55 inputs and 27x27 outputs, and a third convolutional layer with 55x55 inputs and 48x48 outputs. The second set consists of three convolutional layers with 3x3 kernels and 13x13 outputs, followed by two max pooling layers with 13x13 inputs and 192x192 outputs, and a final convolutional layer with 192x192 inputs and 2048x2048 outputs. The final output is a dense layer with 1000 units. The diagram also shows a GPU icon.

of transistors 10^9

of pixels used in training 10^{14}

Fei-Fei Li & Andrej Karpathy

Lecture 1 - 33

5-Jan-15

Exportar archivo PDF

Adobe Export PDF

Convertir archivos PDF a Word o Excel Online

Seleccionar archivo PDF

lecture1.pdf

Convertir a

Microsoft Word (*.docx)

Idioma del documento:

Español Cambiar

Convertir

Crear archivo PDF

Editar PDF

Comentar

Rellenar y firmar

Enviar para firmar

Enviar y realizar un seguimiento

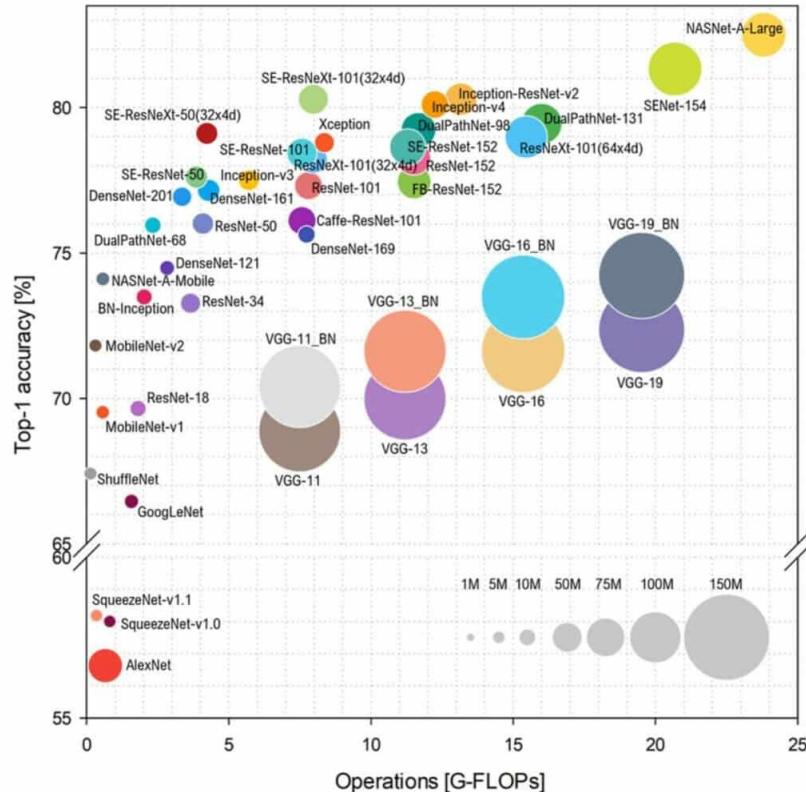
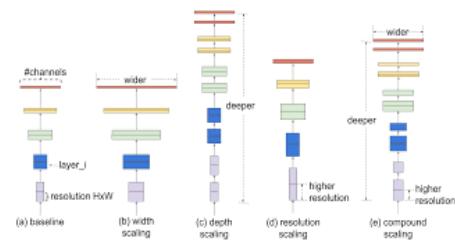
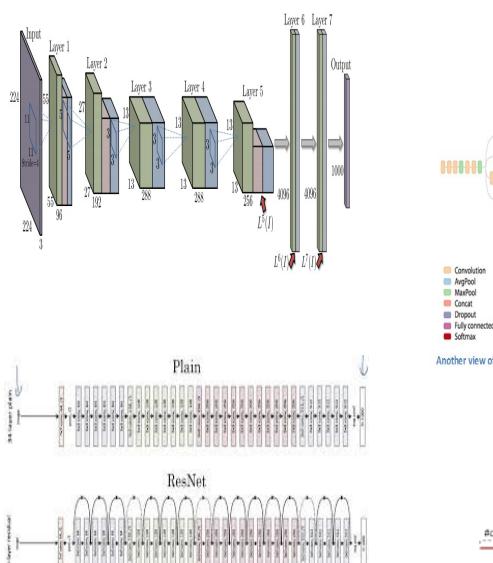
Almacene y comparta archivos en Document Cloud

Más información

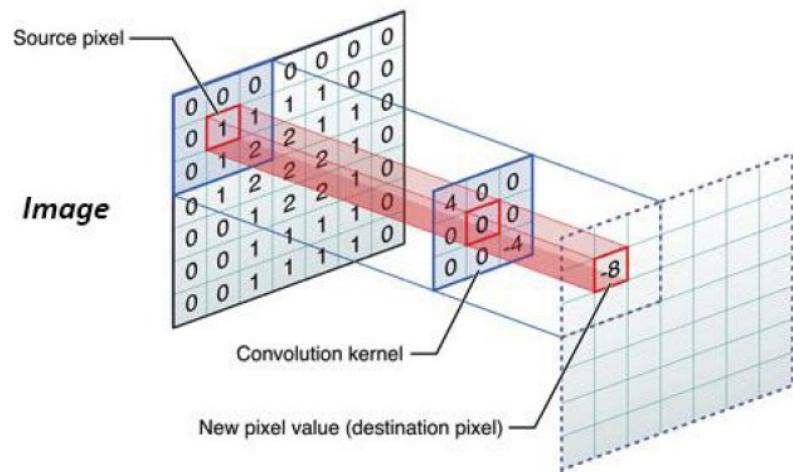
11:33 19/11/2015

NN models

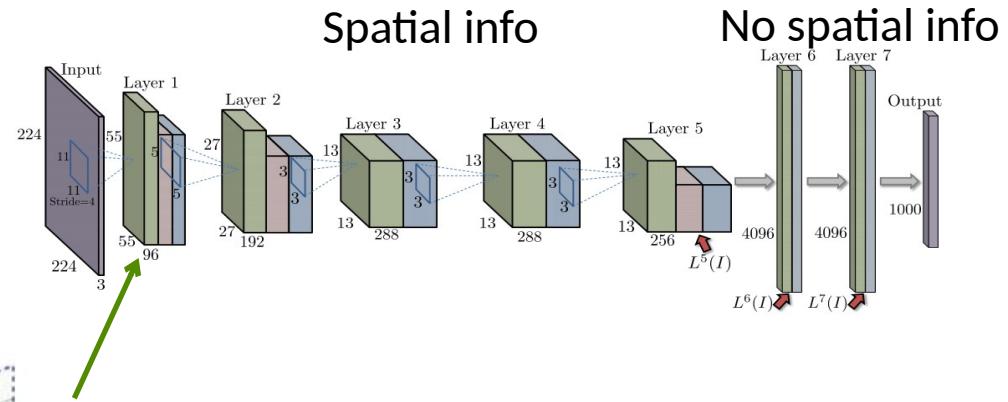
))))



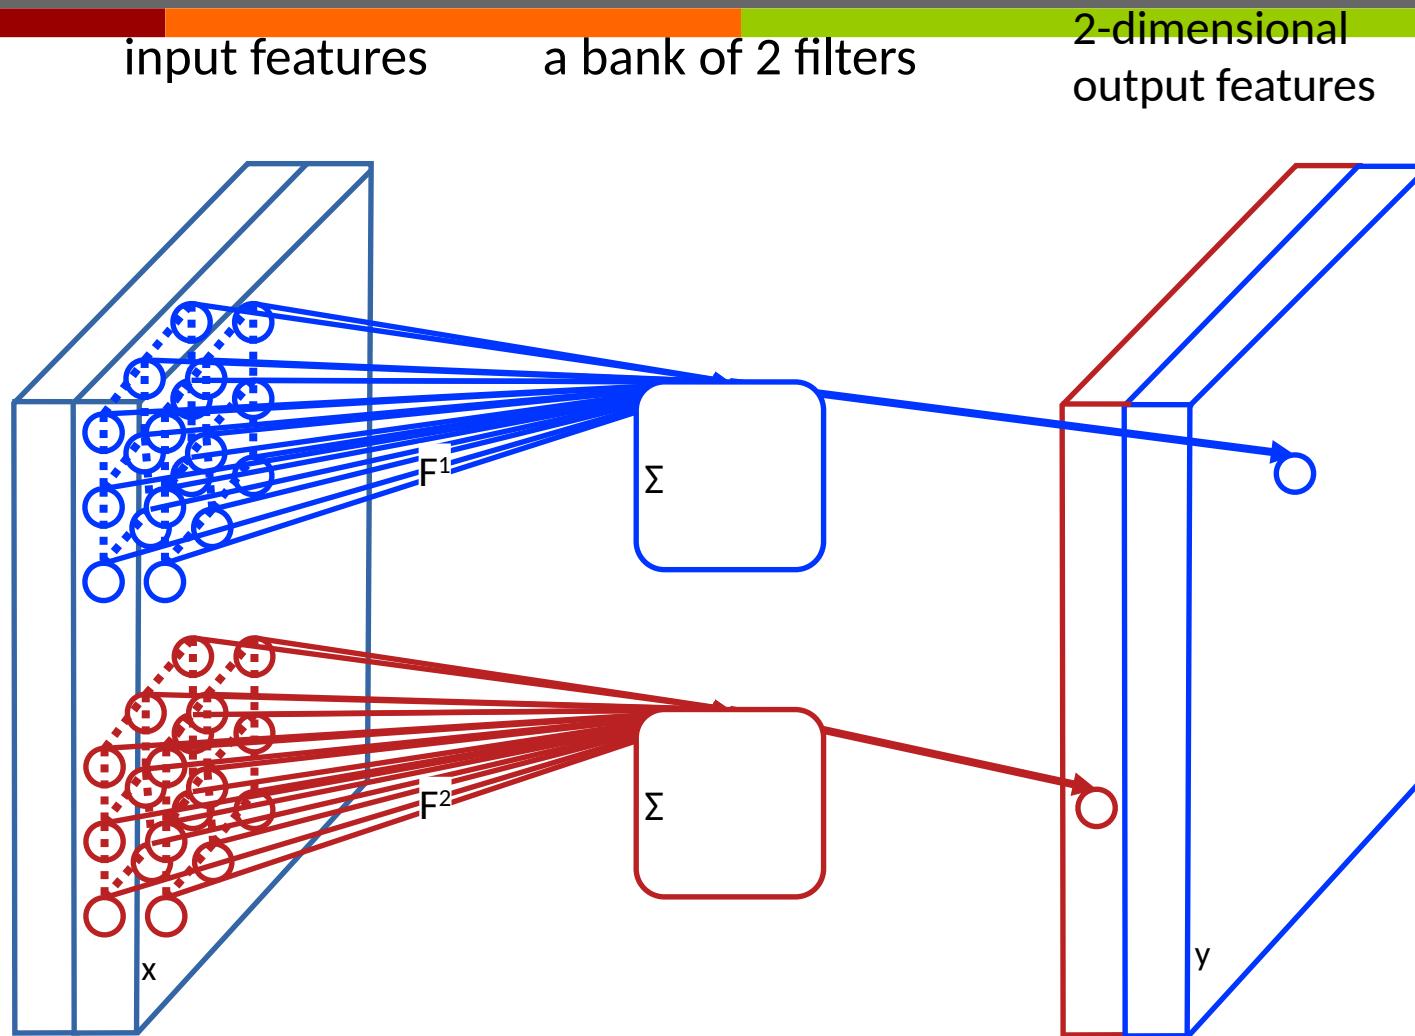
Convolutional and Max-pooling layer



Convolutional layer

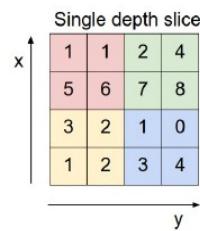
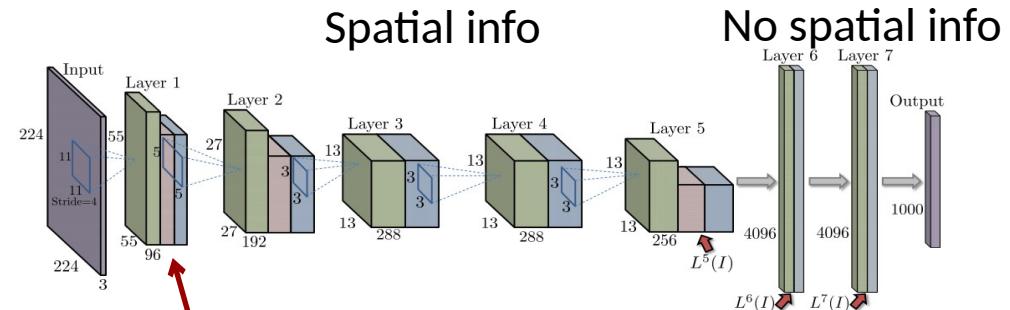


Why is it convolutional?

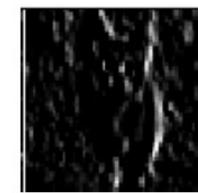


Convolution

Convolutional and Max-pooling layer

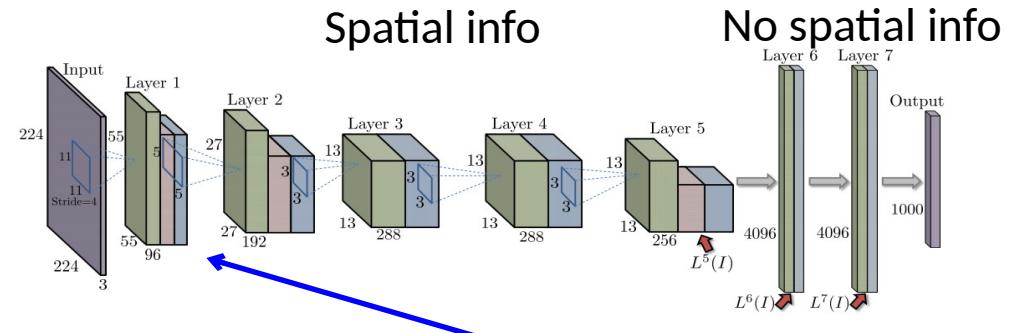


max pool with 2x2 filters
and stride 2



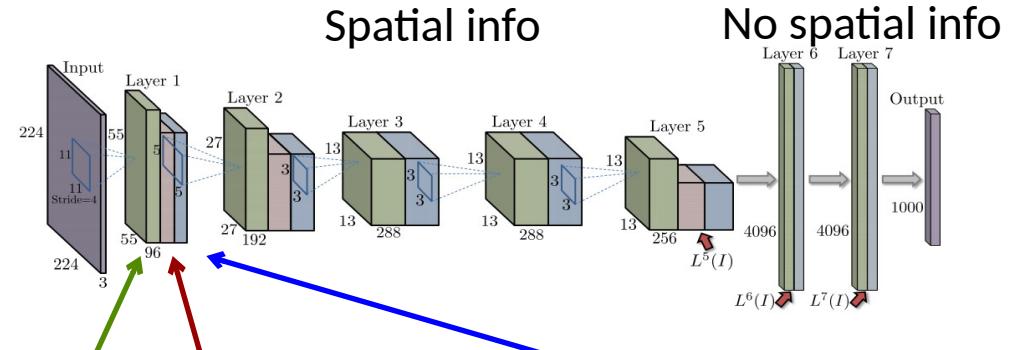
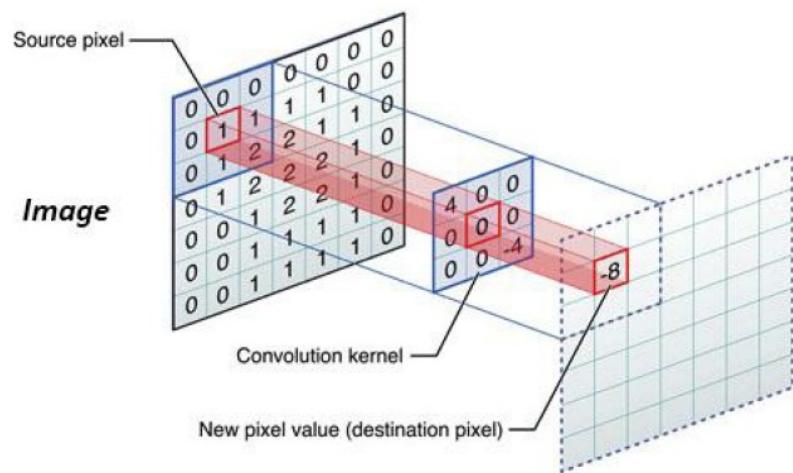
Max-pool layer

Convolutional and Max-pooling layer



ReLU $\max(0, x)$

Convolutional and Max-pooling layer



ReLU $\max(0, x)$

Convolutional layer

Single depth slice			
x	1	1	2
y	5	6	7
1	2	1	0
1	2	3	4

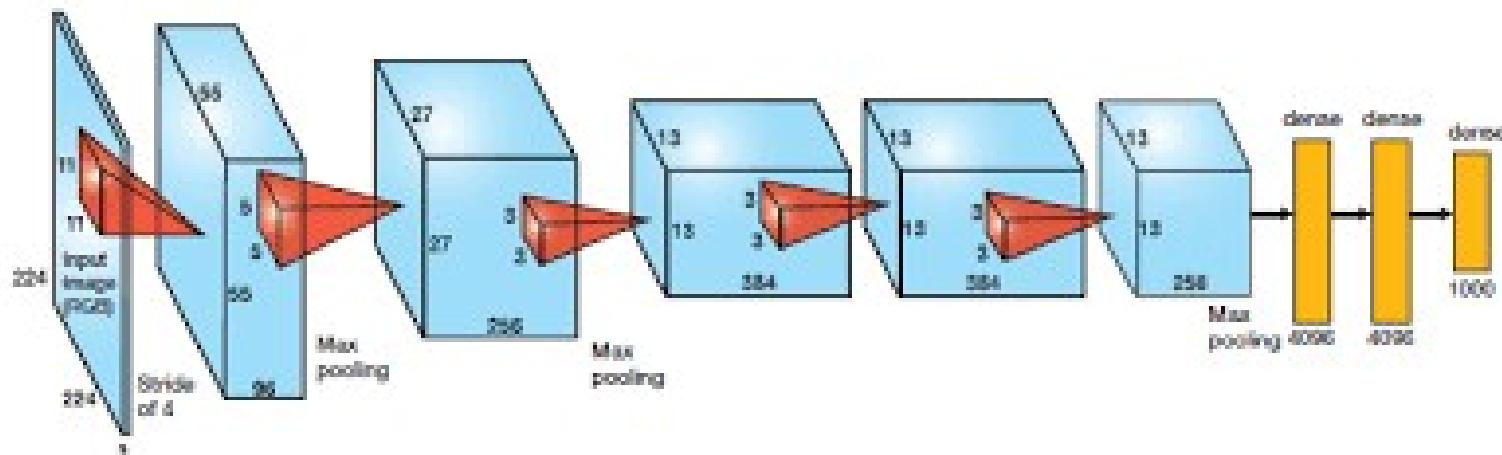
max pool with 2x2 filters and stride 2

6	8
3	4

Max-pool layer

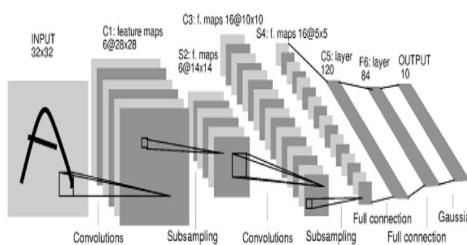


Why is it convolutional?



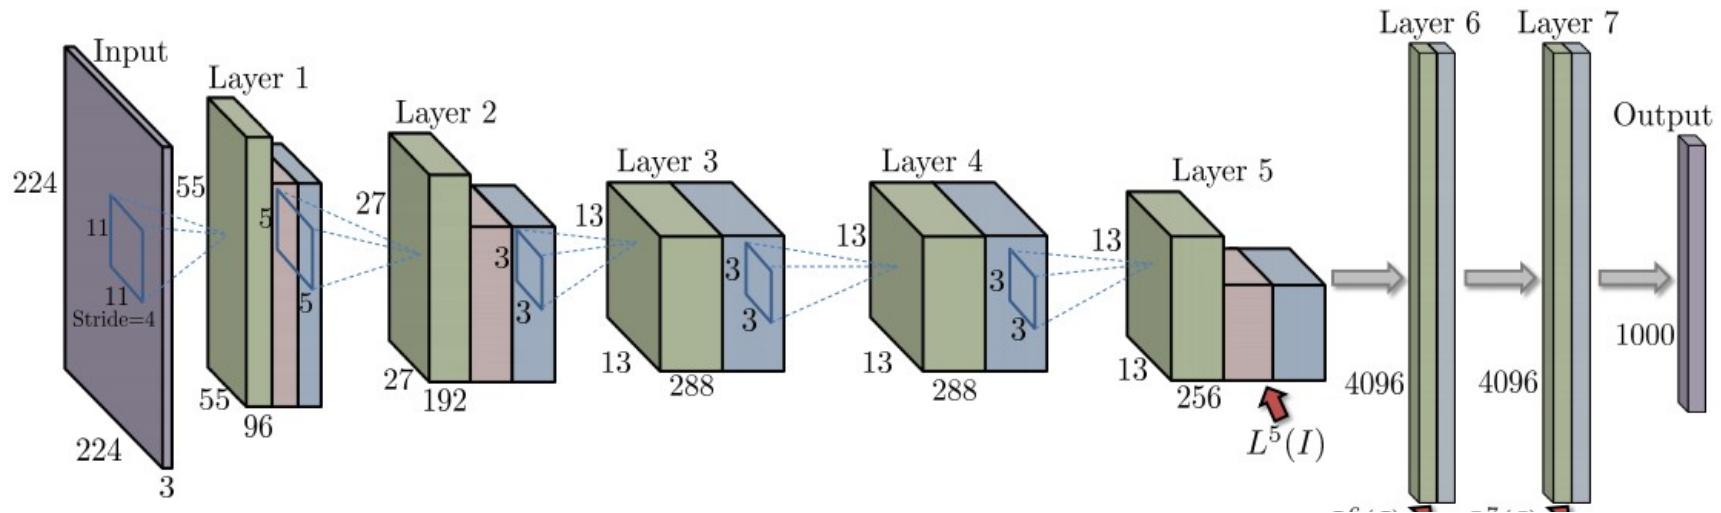
1998

LeCun et al.



LeCun, Chief AI Scientist for Facebook AI Research (FAIR), and a Silver Professor at New York University

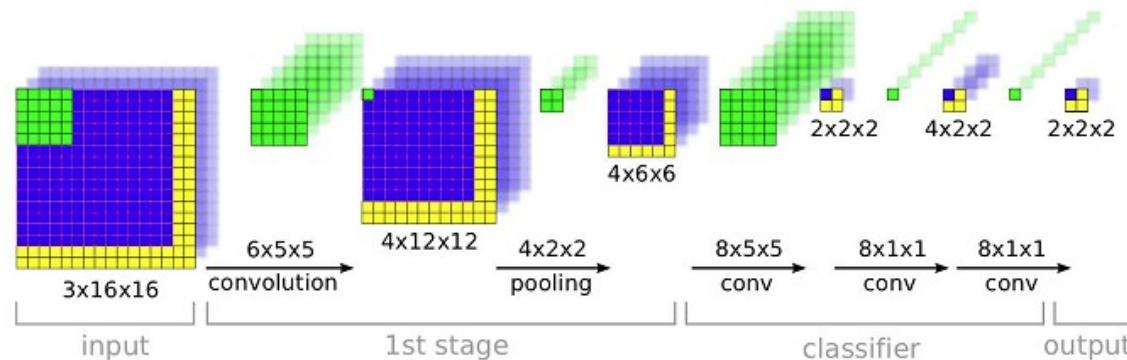
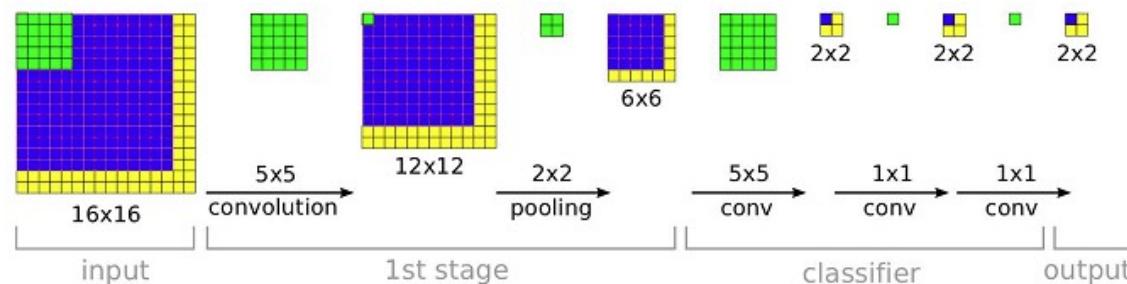
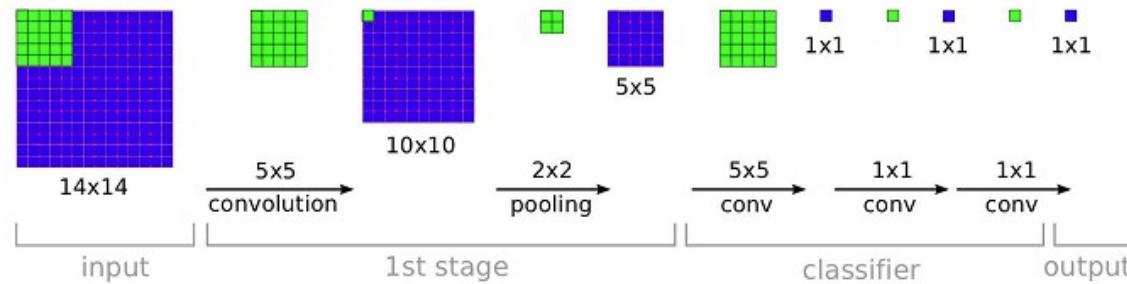
Training a CNN



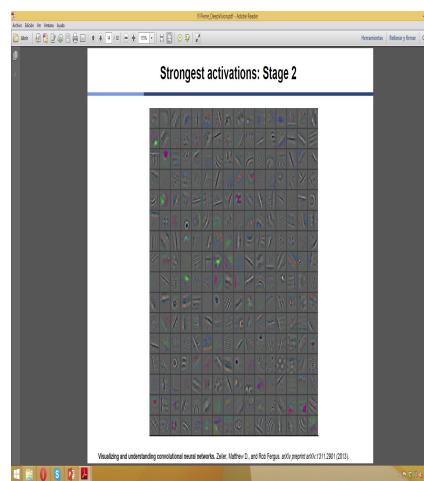
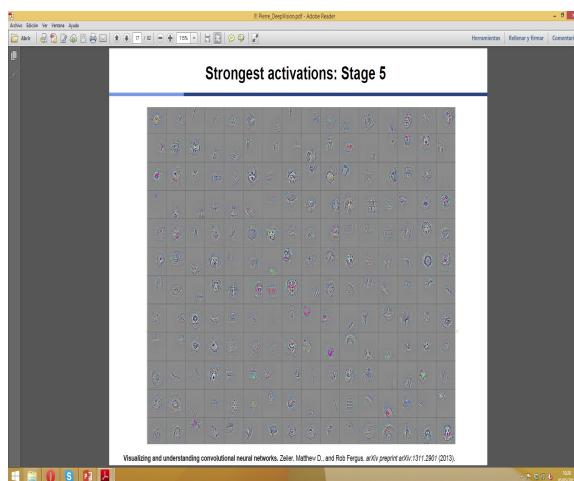
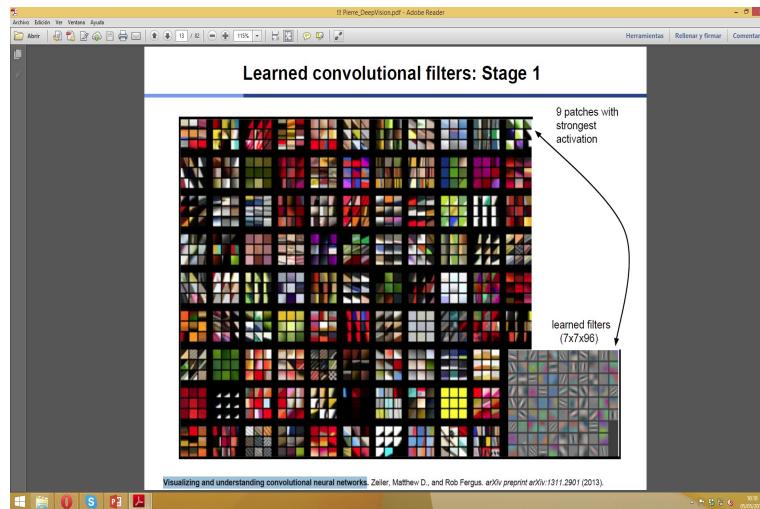
The process of training a CNN consists of training all hyperparameters: convolutional matrices and weights of the fully connected layers.

- **Several millions of parameters!!!**

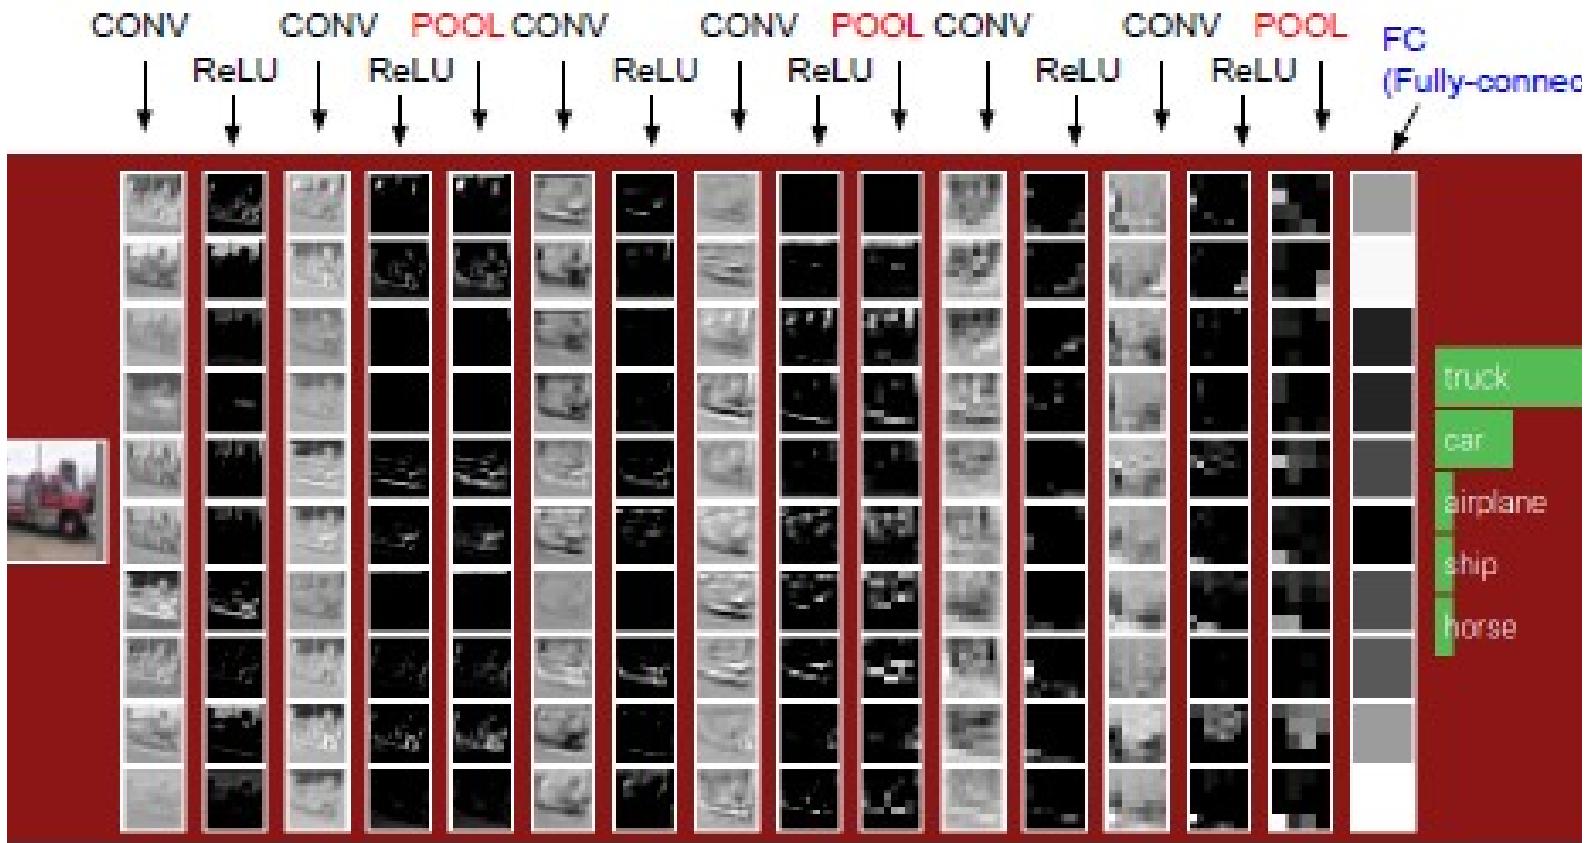
How does the CNN work?



Learned convolutional filters



Example architecture



The trick is to train the weights such that when the network sees a picture of a truck, the last layer will say “truck”.

1001 benefits of CNN

|

Transfer learning: Fine tuning for object recognition

Replace and retrain the classifier on top of the ConvNet

Fine-tune the weights of the pre-trained network by continuing the backpropagation

Feature extraction by CNN

Object detection

Object segmentation

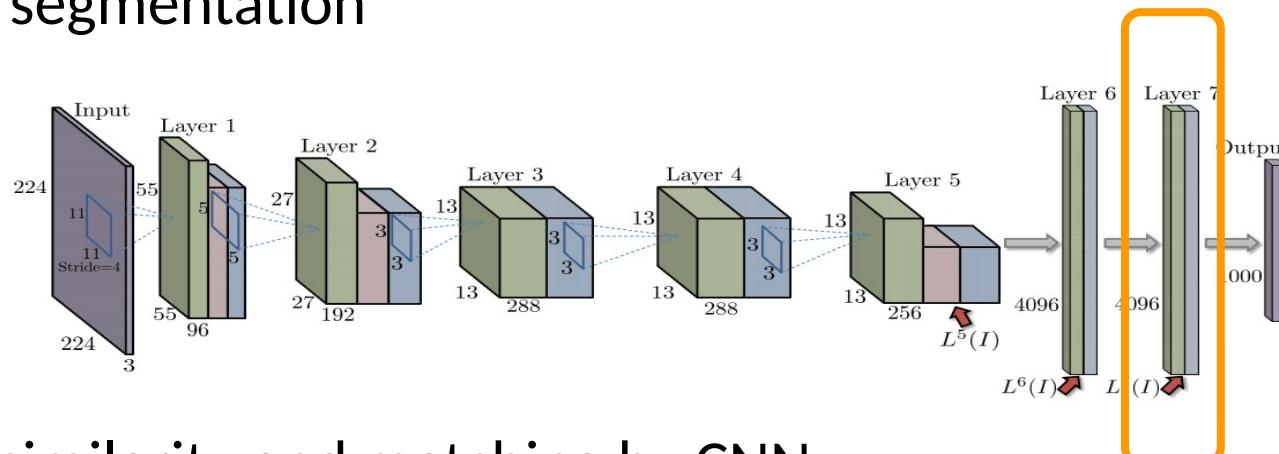
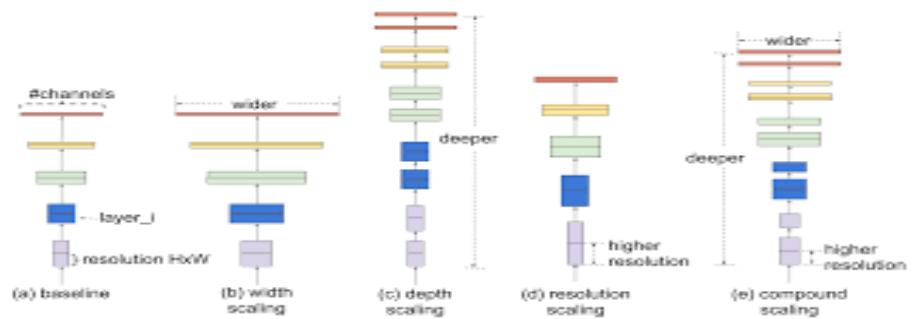
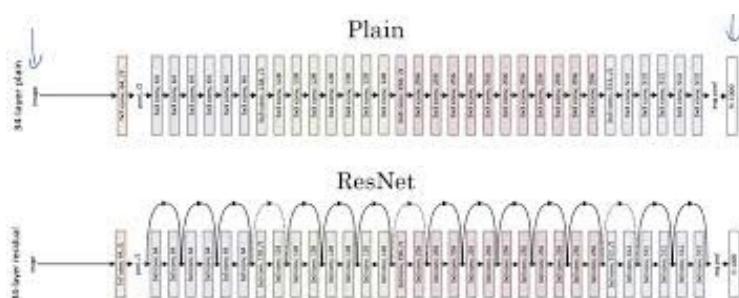
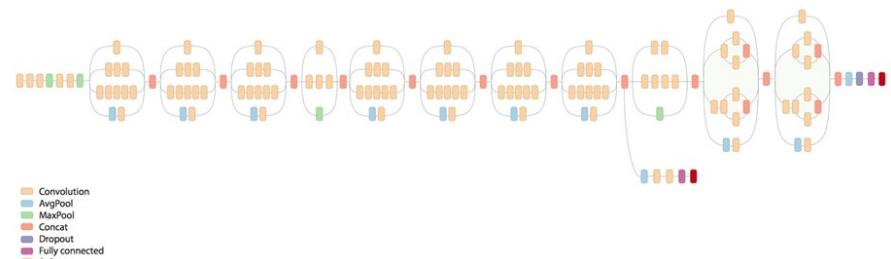
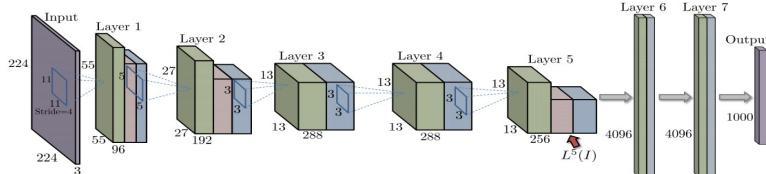


Image similarity and matching by CNN

Convolutional Neural Networks (4096 Features)

CNN models

))))



IMAGENET Large Scale Visual Recognition Challenge

Year 2010

NEC-UIUC



Dense descriptor grid:
HOG, LBP

Coding: local coordinate,
super-vector

Pooling, SPM

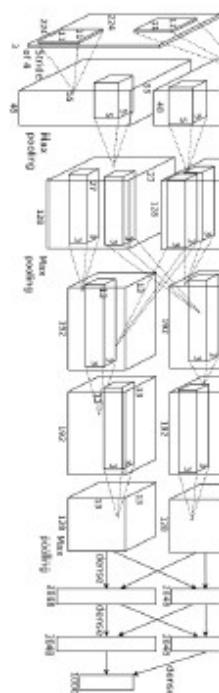
Linear SVM

[Lin CVPR 2011]

Lion image by Swissfrog
is
licensed under CC BY 3.0

Year 2012

SuperVision



[Krizhevsky NIPS 2012]

Figure copyright Alex Krizhevsky, Ilya
Sutskever, and Geoffrey Hinton, 2012.
Reproduced with permission.

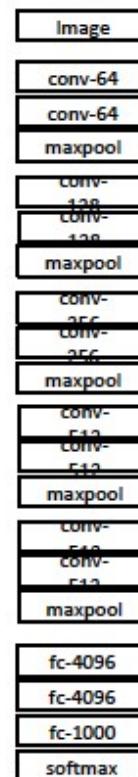
Year 2014

GoogLeNet



[Szegedy arxiv 2014]

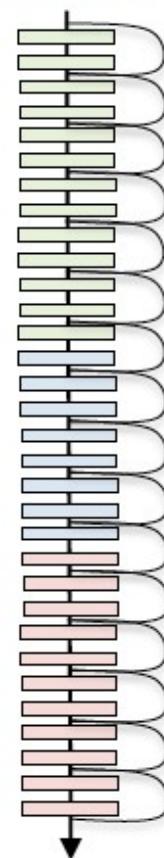
VGG



[Simonyan arxiv 2014]

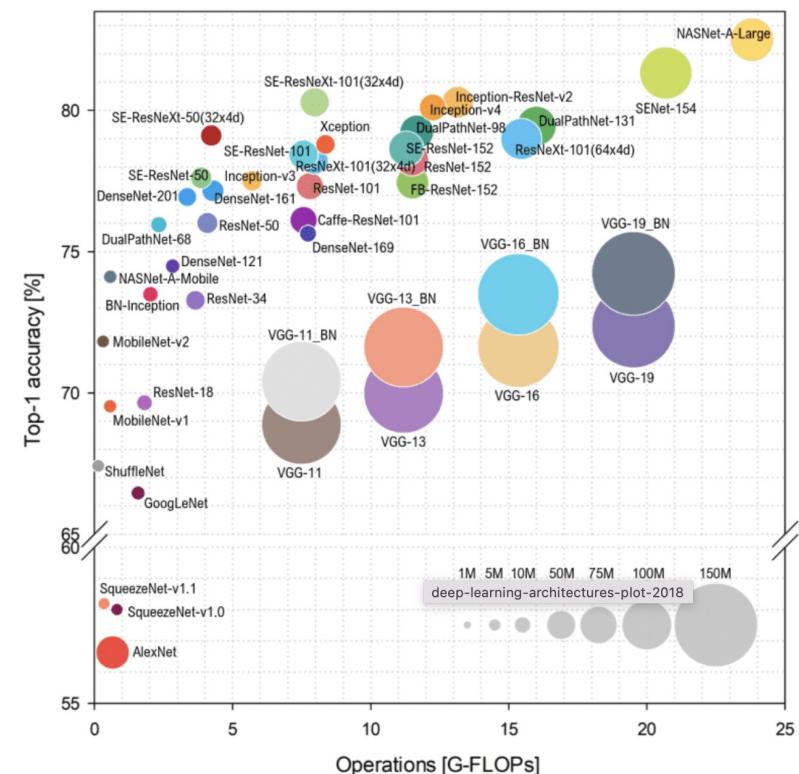
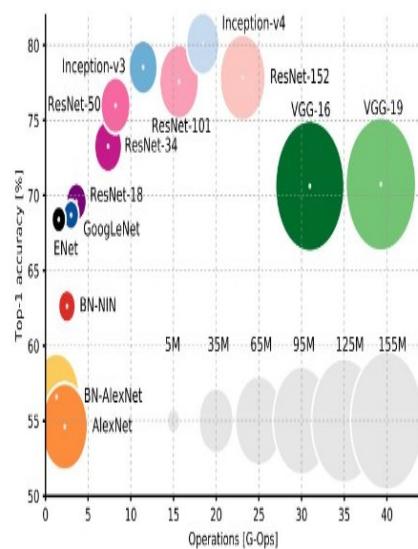
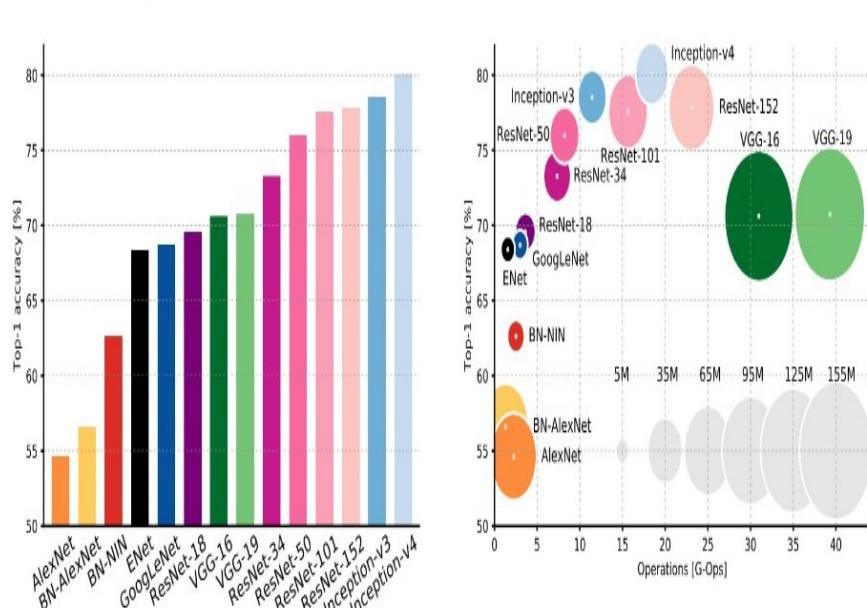
Year 2015

MSRA



[He ICCV 2015]

Analysis of CNNs



An Analysis of Deep Neural Network Models for Practical Applications, 2017.

- Millions of parameters!!!

The process of training a CNN consists of training all hyperparameters: convolutional matrices and weights of the fully connected layers.

- ↗ AI, Machine learning & Deep learning
- ↗ What is a Convolutional Neural Network?
 - ↗ Layers
 - ↗ Loss function and CNN Optimization
- ↗ Applications

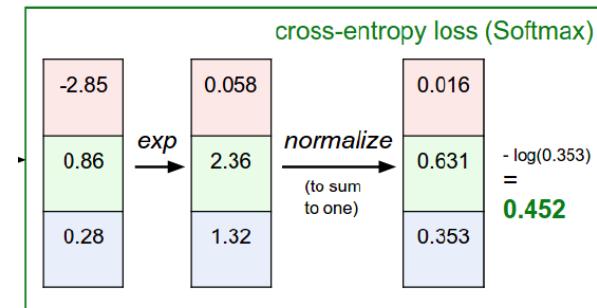
Loss function and optimisation



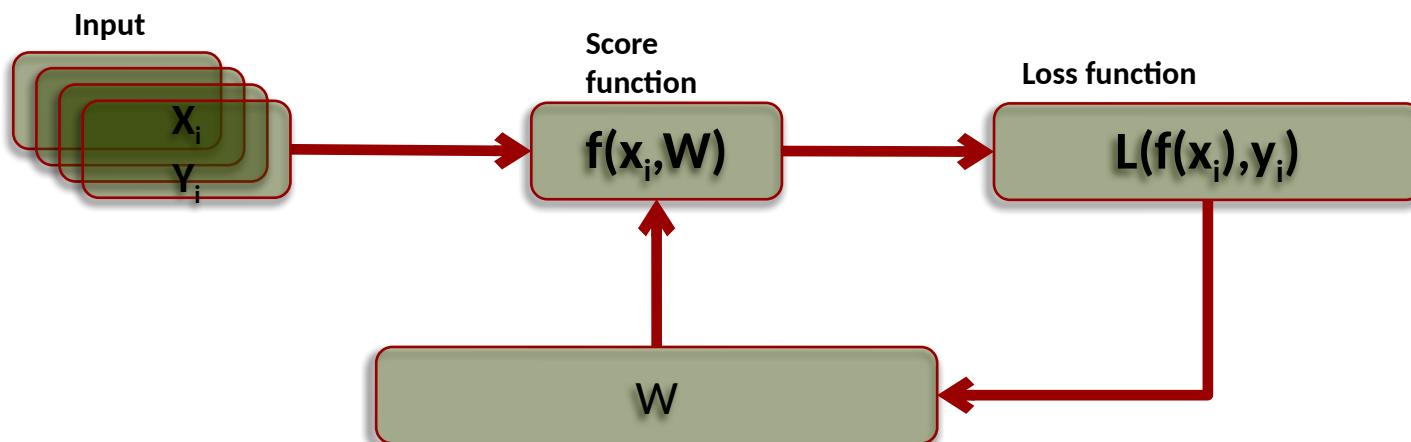
Question: if you were to assign a single number to how unhappy you are with these scores, what would you do?

$$L_i = -\log \left(\frac{e^{f_{y_i}}}{\sum_j e^{f_j}} \right)$$

softmax function



Question : Given the score and the loss function, how to find the parameters W?



Single-label classification

- Change the lost function to the Binary cross-entropy function L_b :

Conventional approach:
CNN for classification



Softmax

$$P(y_i|x) = \frac{\exp^{f(x)_i}}{\sum_i \exp^{f(x)_i}}$$

loss

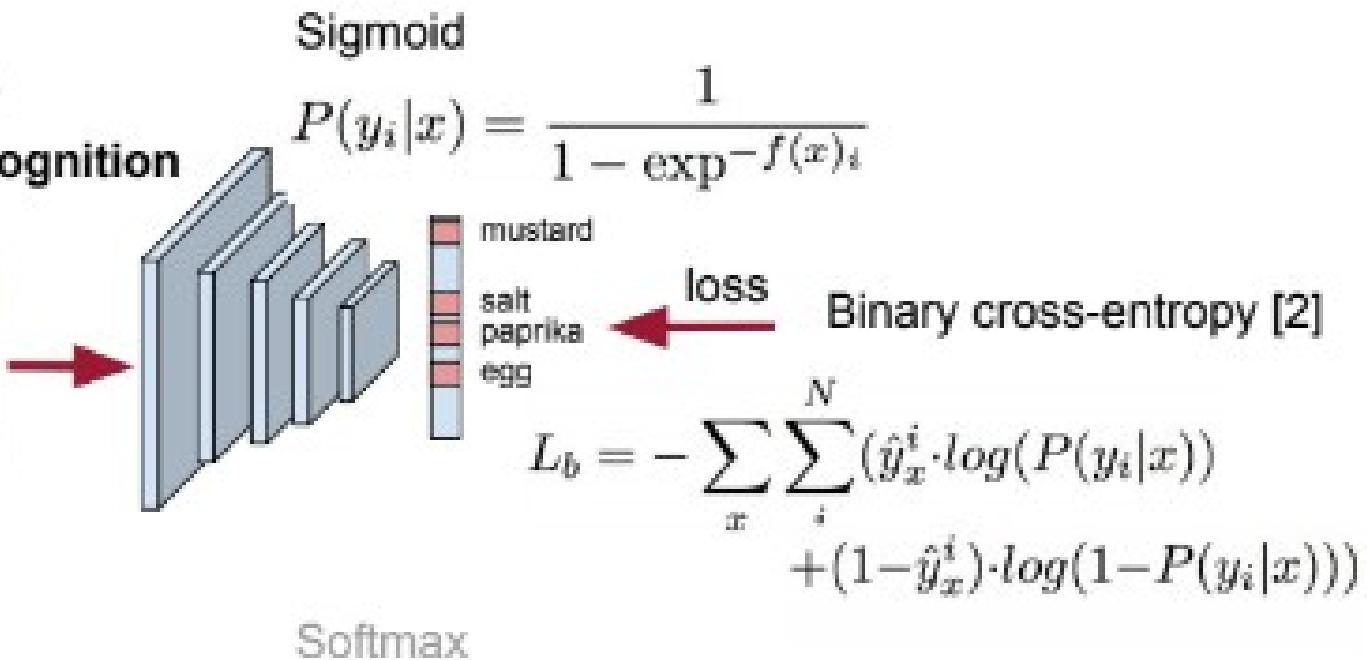
Categorical cross-entropy

$$L_c = - \sum_x \log(P(\hat{y}_x|x))$$

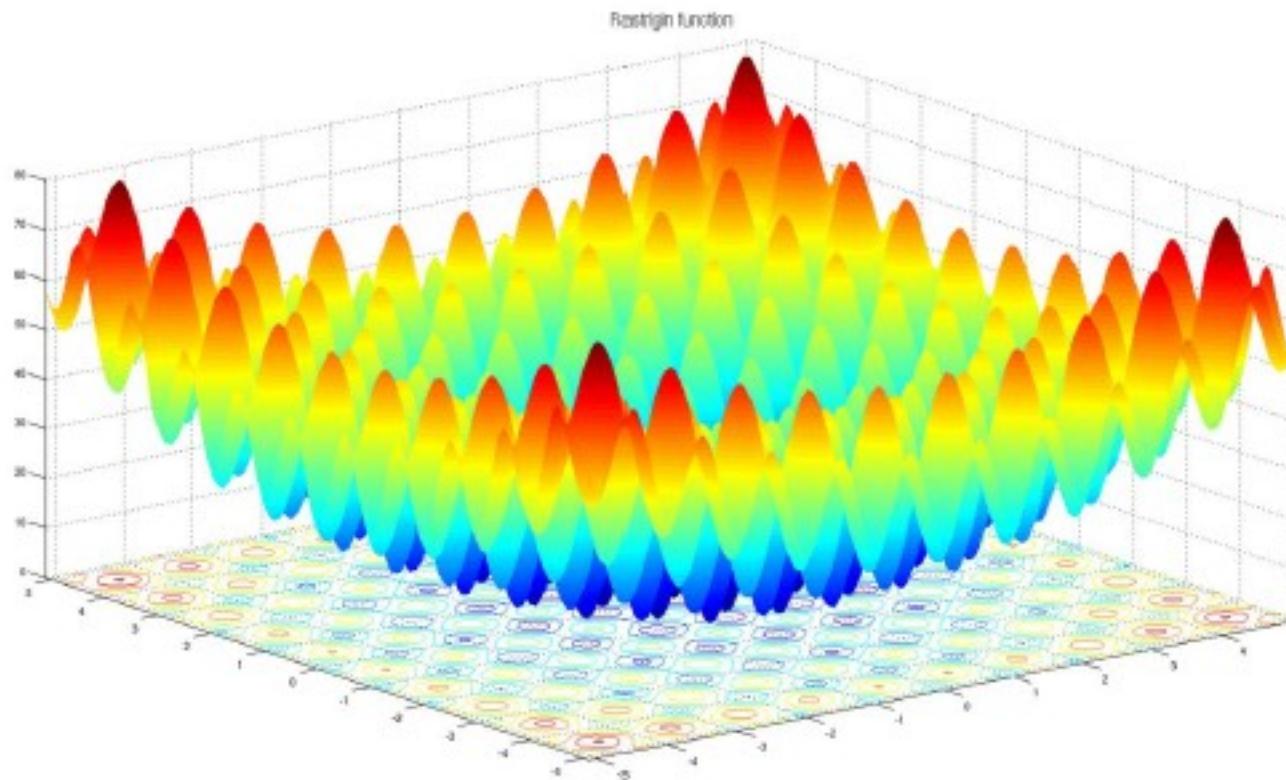
Multi-label classification

- ↗ Change the lost function to the Binary cross-entropy function L_b :

**Our proposal:
Adaptation for
multi-label recognition**



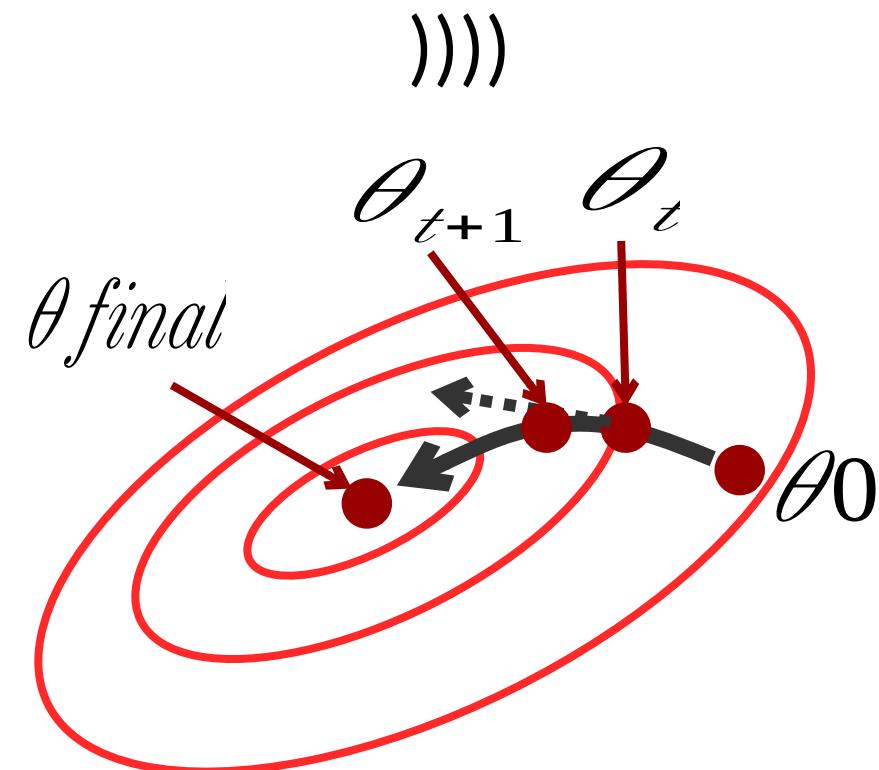
Optimization



The process of training a CNN consists of training all hyperparameters: convolutional matrices and weights of the fully connected layers.

- Millions of parameters!!!

Gradient descent



- ↗ Initialize randomly
- ↗ For t in $0, \dots, T_{\text{maxiter}}$
 - ↗ Gradient of the objective
 - ↗ Learning rate (step size)

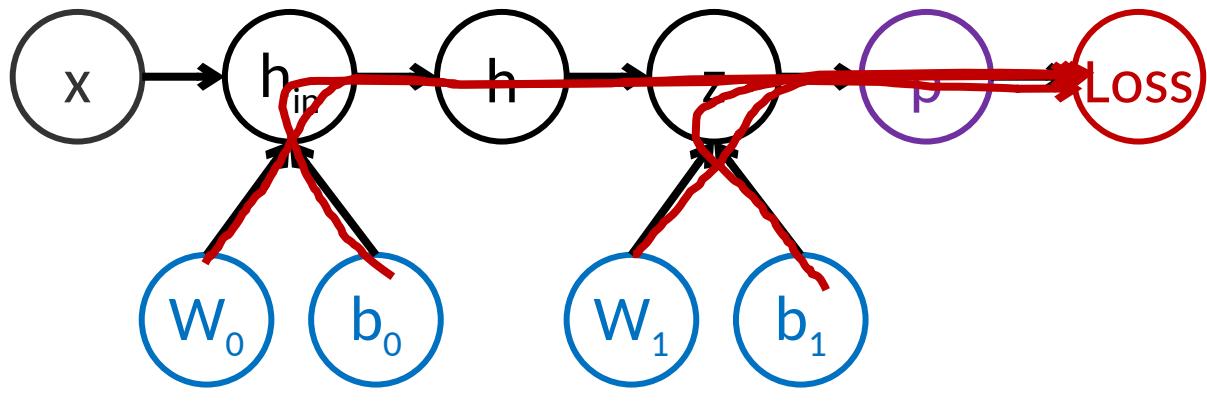
- Computation of $J(\theta)$ requires a full sweep over the training data
- Per-iteration comp. cost = $O(n)$

Chain rule

- Identify how each variable influence the loss

$$\frac{\partial \text{Loss}}{\partial W_1} = \frac{\partial \text{Loss}}{\partial p} \cdot \dots \cdot \frac{\partial z}{\partial W_1}$$

$$\frac{\partial \text{Loss}}{\partial b_1} = \frac{\partial \text{Loss}}{\partial p} \cdot \dots \cdot \frac{\partial z}{\partial b_1}$$

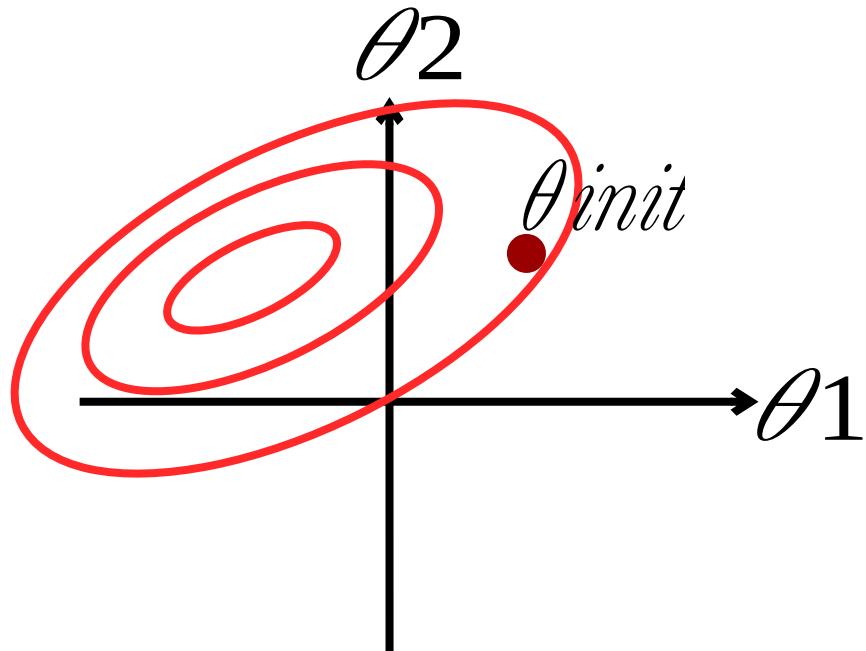


$$\frac{\partial \text{Loss}}{\partial W_0} = \frac{\partial \text{Loss}}{\partial p} \cdot \dots \cdot \frac{\partial h_{in}}{\partial W_0}$$

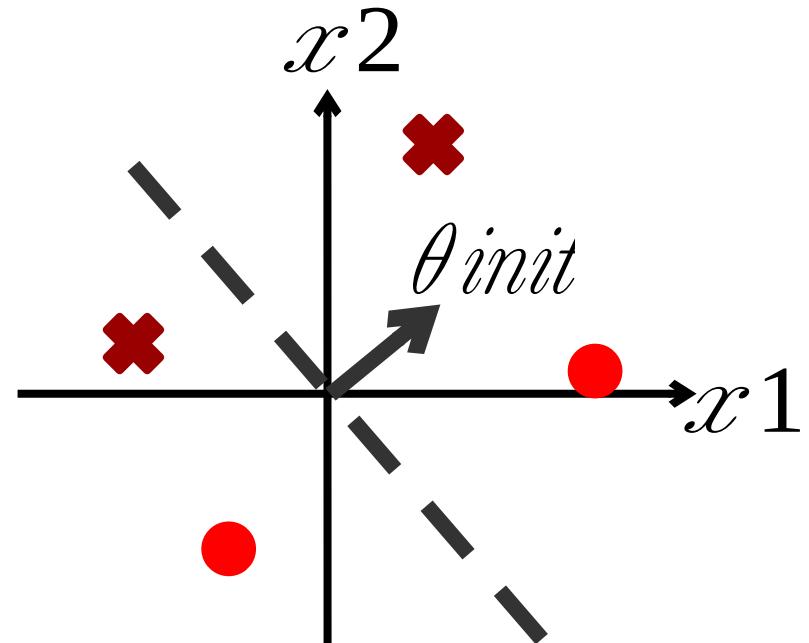
$$\frac{\partial \text{Loss}}{\partial b_0} = \frac{\partial \text{Loss}}{\partial p} \cdot \dots \cdot \frac{\partial h_{in}}{\partial b_0}$$

Landscape of training objective

Parameter space

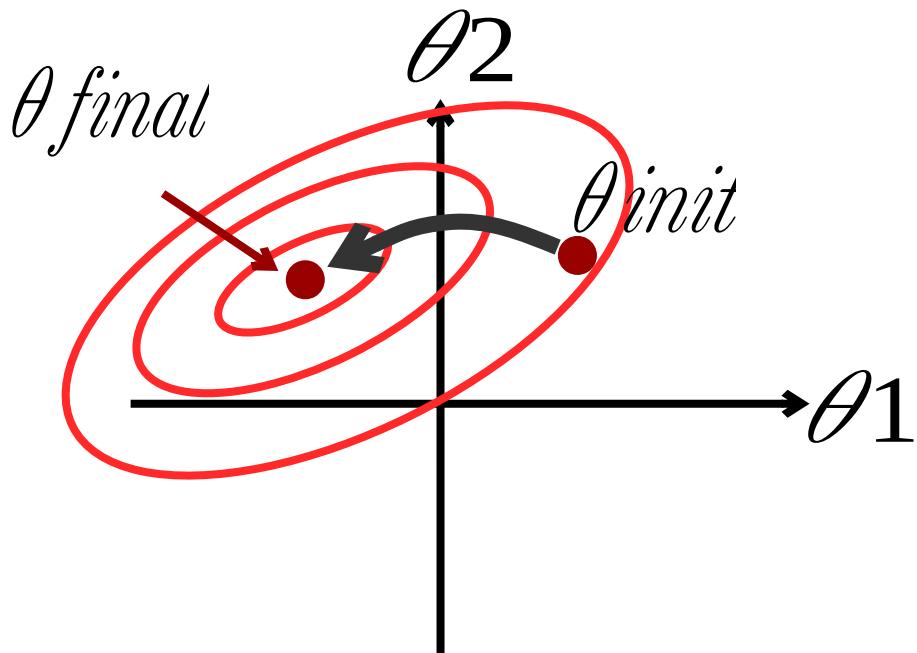


Example space

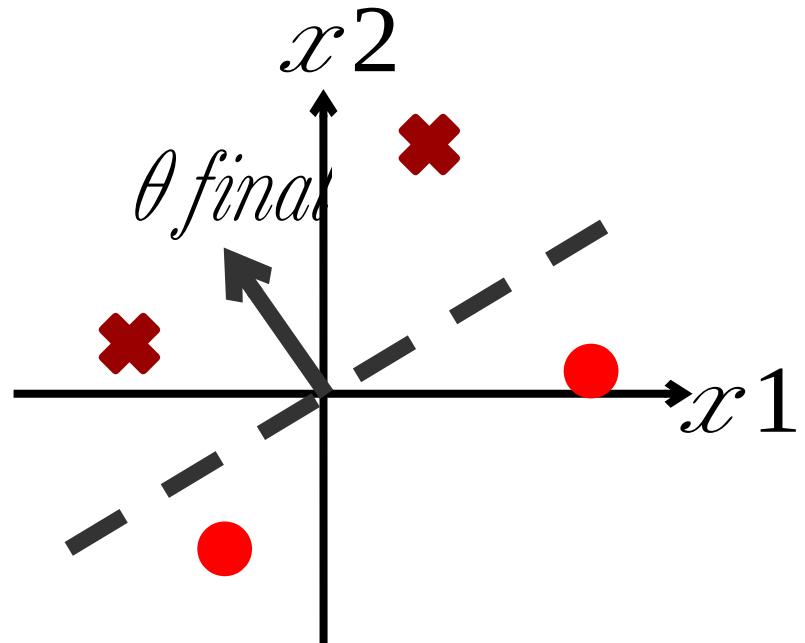


Landscape of training objective

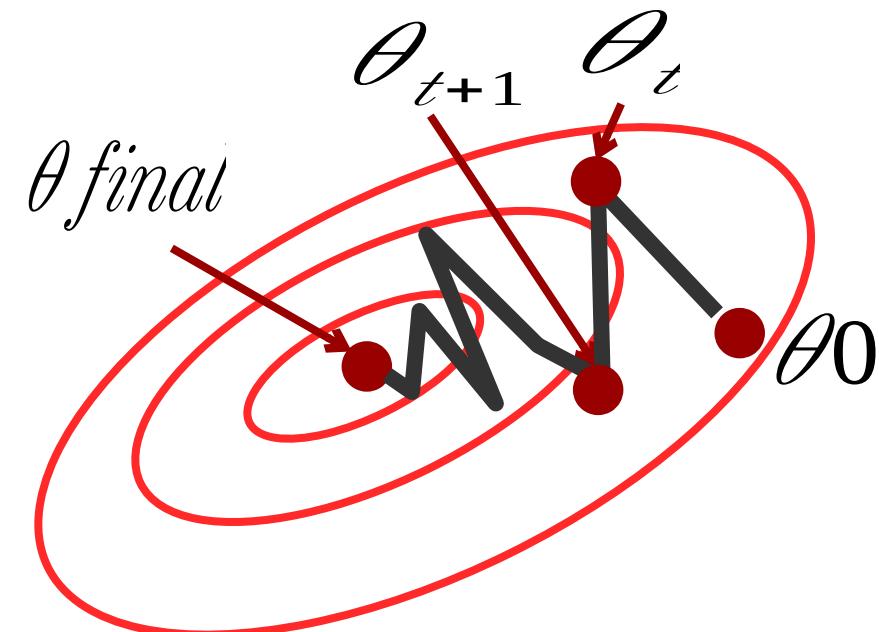
Parameter space



Example space



Stochastic gradient descent (SGD)



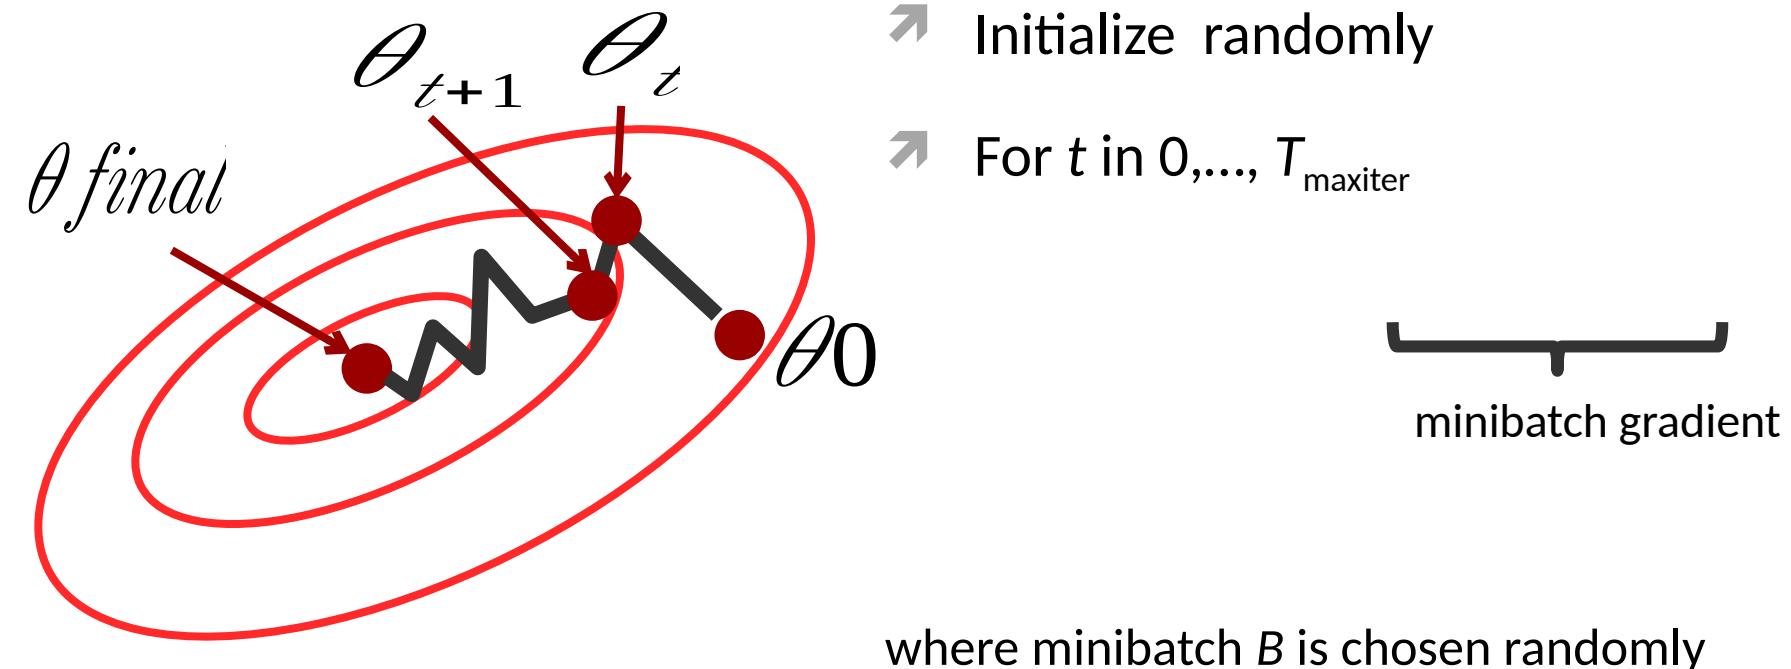
- ↗ Initialize randomly
- ↗ For t in $0, \dots, T_{\text{maxiter}}$

Stochastic gradient

where index i is chosen randomly

- computation of $J(\theta)$ requires only one training example
- Per-iteration comp. cost = $O(1)$

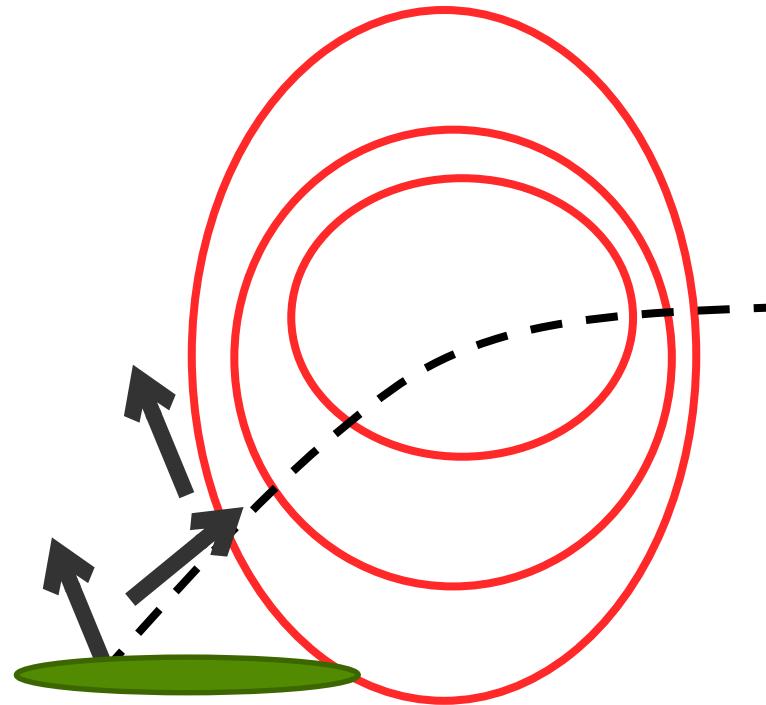
Mini-batch stochastic gradient descent



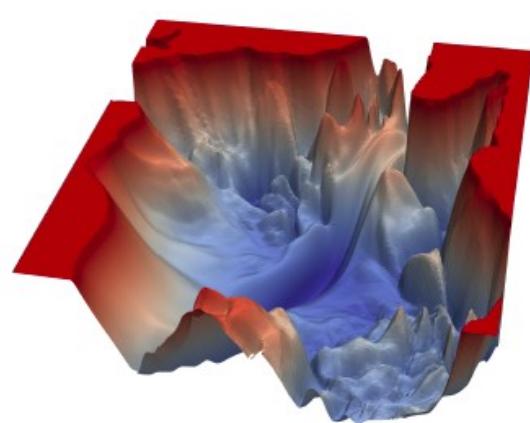
- is average gradient over random subset of data of size B
- Per-iteration comp. cost = $O(B)$

More optimization algorithms

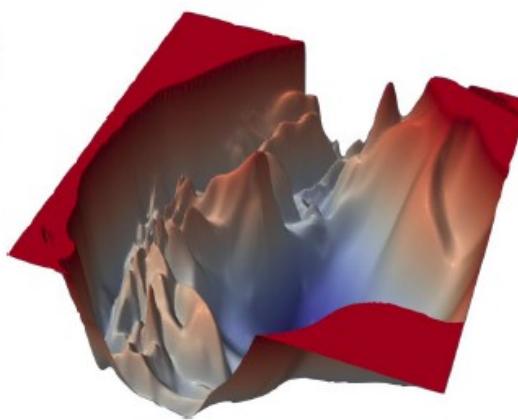
- ↗ **Momentum SGD:** improves SGD by incorporating “momentum”
- ↗ **Adam** [Kingma & Ba 2015]: uses first and second order statistics of the gradients so that gradients are normalized
- ↗ **Benefit:** prevents the vanishing/exploding gradient problem



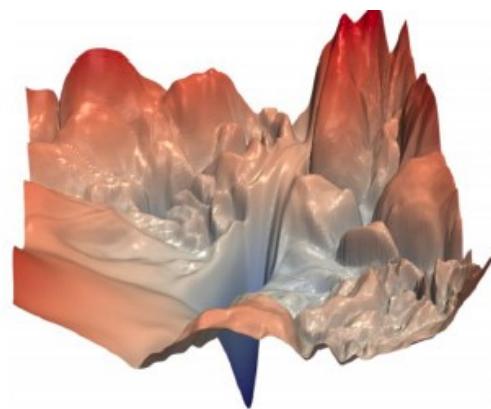
Stochastic Gradient Descent



VGG-56



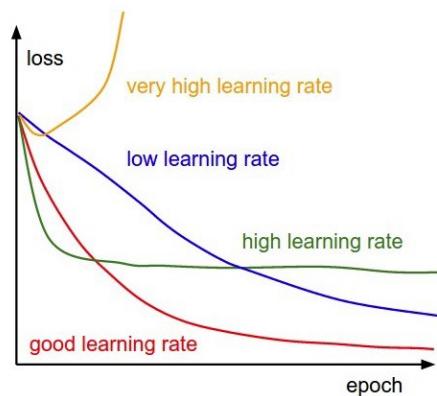
VGG-110



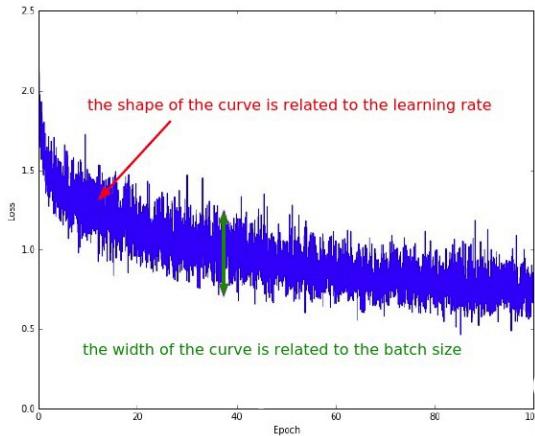
ResNet-56

Hao Li et al., NIPS, 2017

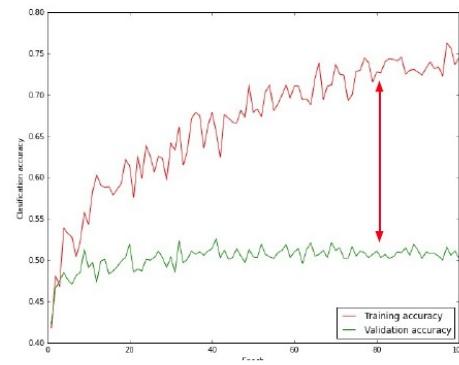
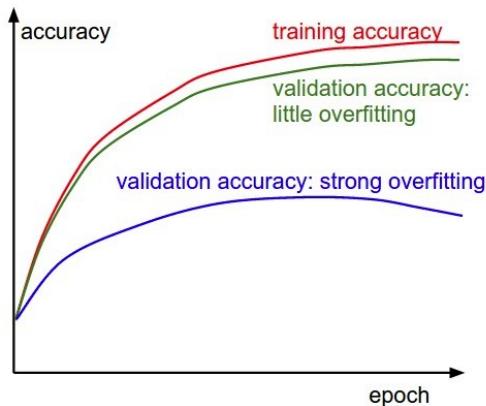
Monitoring loss and accuracy



Adapted from Fei Fei slides



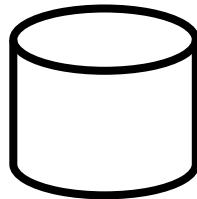
Looks linear?
Learning rate too low!
Decreases too slowly?
Learning rate too high.
Looks too noisy?
Increases the batch size.



Big gap?
- you're overfitting,
increase
regularization!

Overfitting – what is signal vs noise?

↗ Imagine:



Training
data



cat



dog

|-----|

Validation data



- ↗ Powerful models are more likely to overfit
- ↗ We need validation data: leave out some portion of the training data to validate the generalizability of the model

Overfitting



With four parameters I can fit an elephant, and with five I can make him wiggle his trunk.

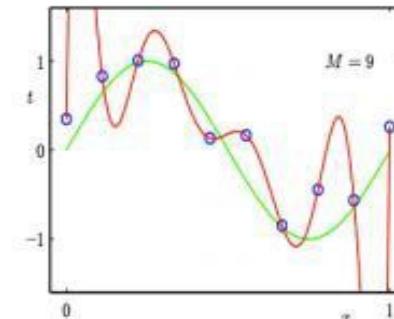
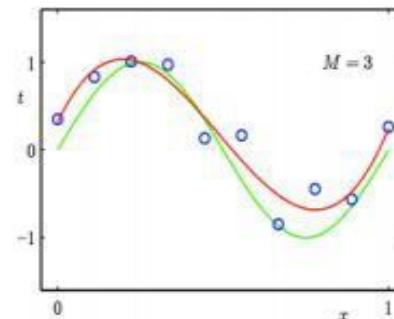
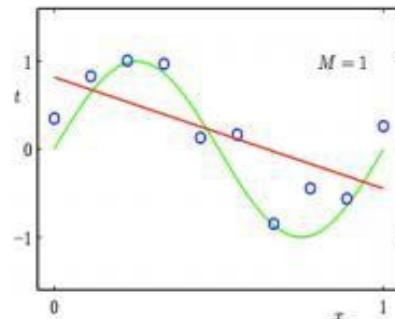
— *John von Neumann* —

AZ QUOTES

Under- and Over-fitting

Under- and Over-fitting examples

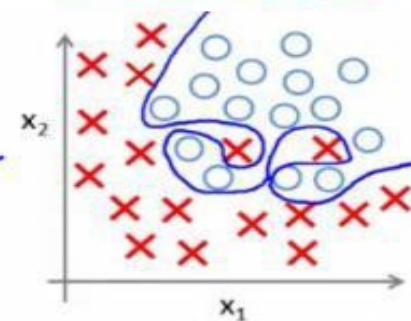
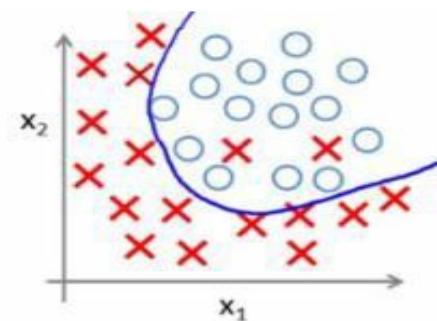
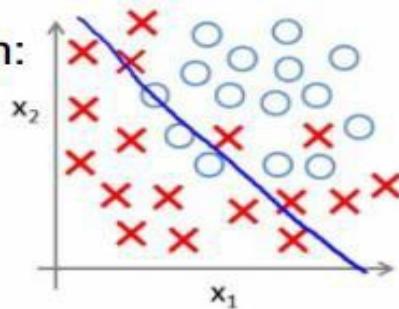
Regression:



predictor too inflexible:
cannot capture pattern

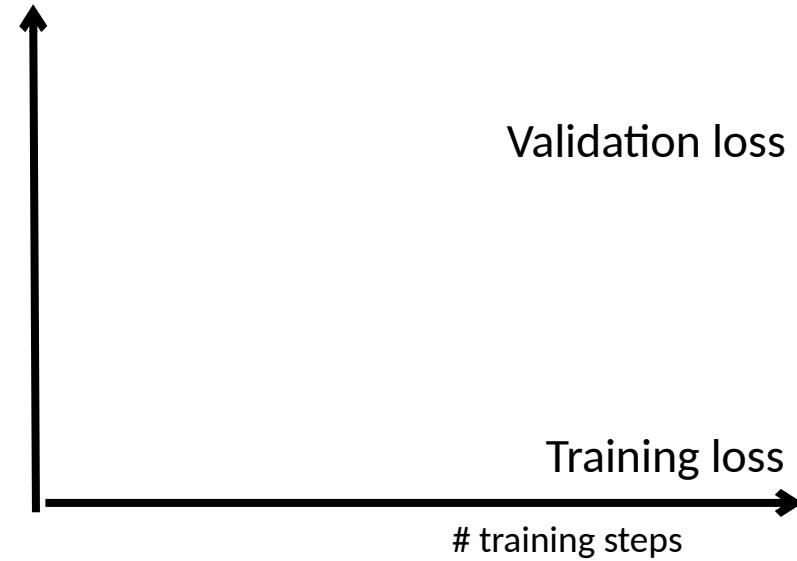
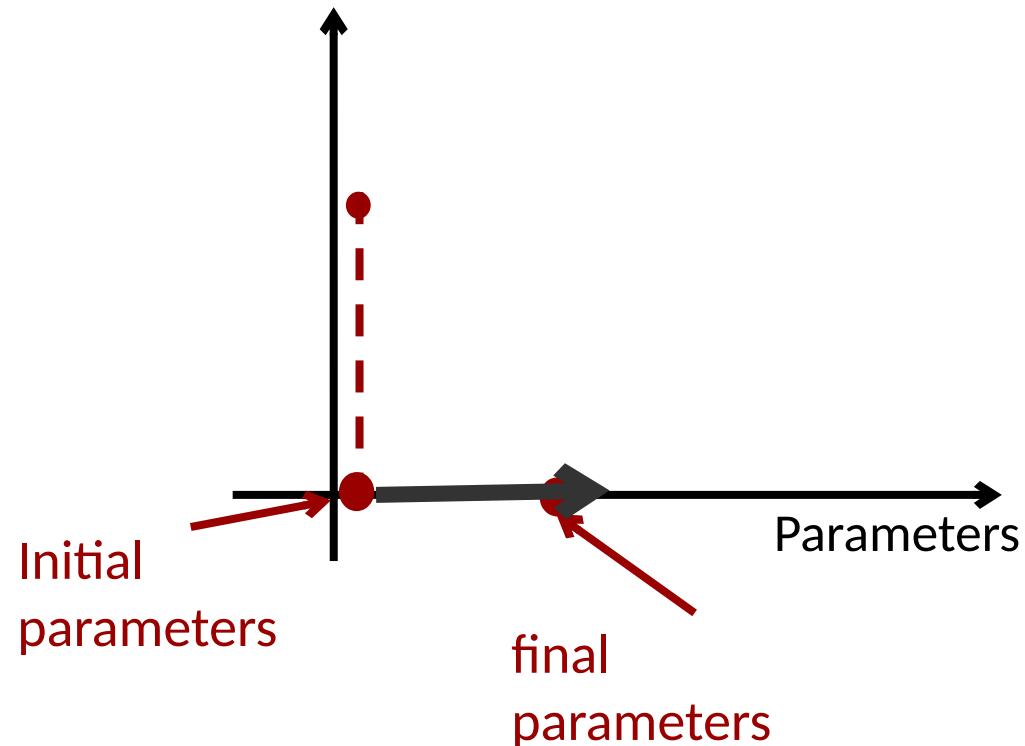
predictor too flexible:
fits noise in the data

Classification:



Landscape of training objective

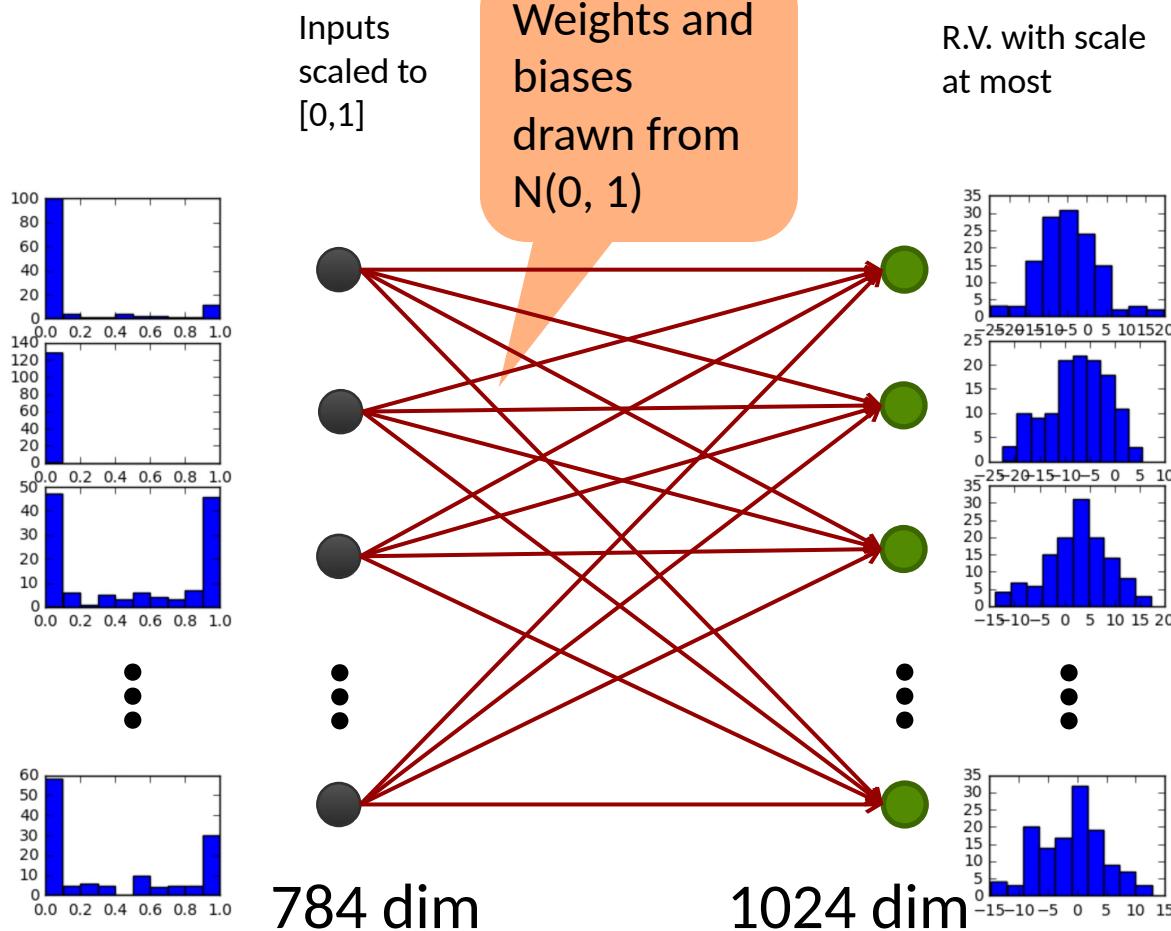
Objective function:



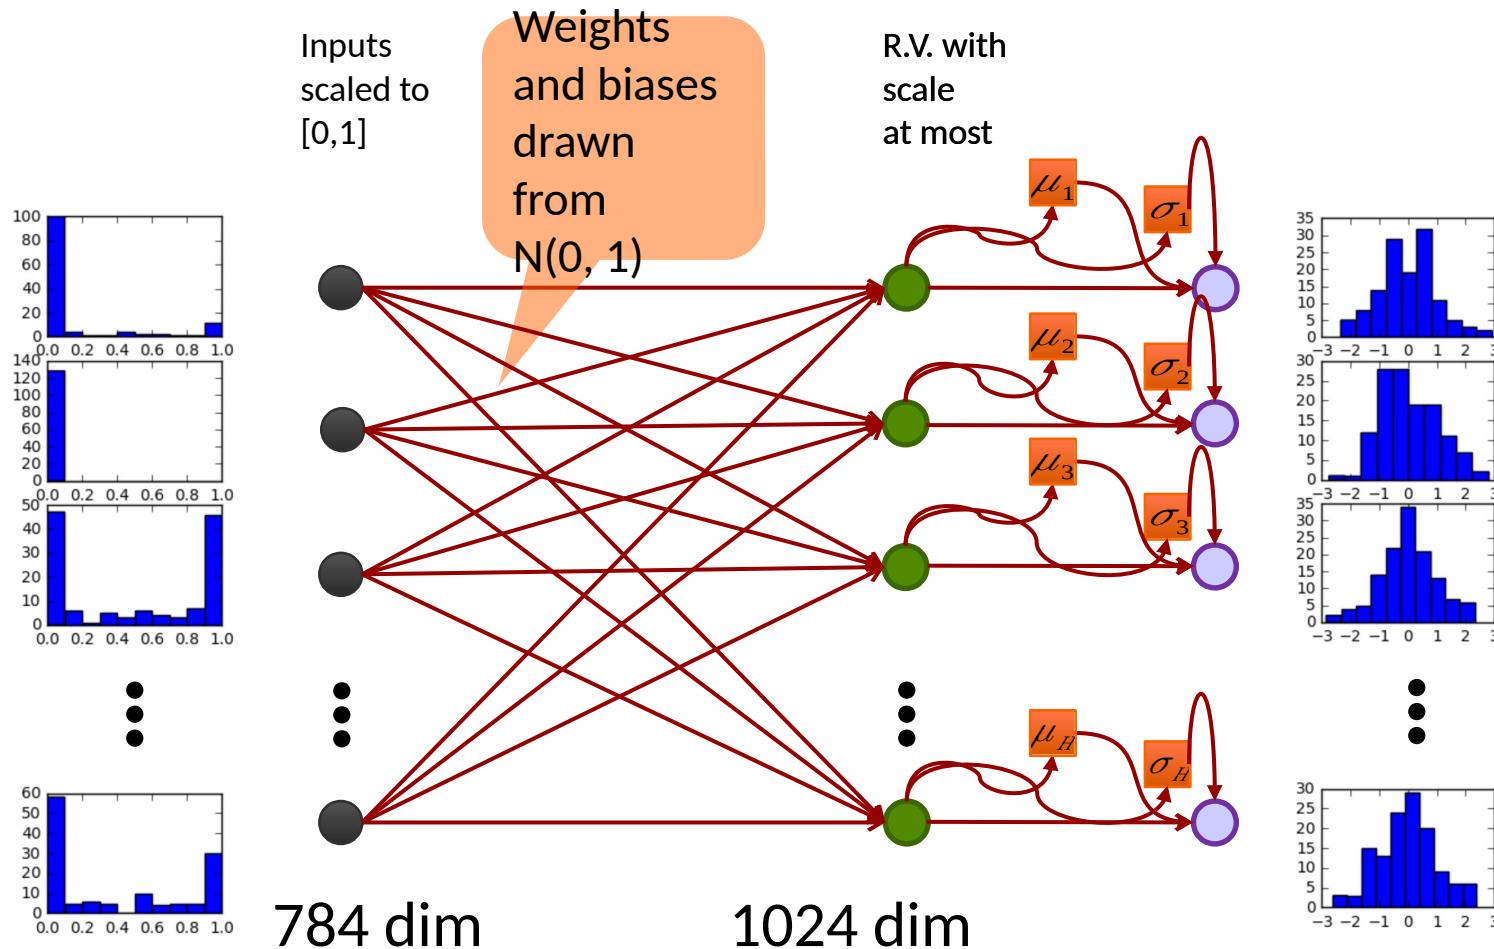
Techniques to reduce overfitting

- ↗ Reduce the number of parameters
 - ↗ Parameter sharing (convnets, recurrent neural nets)
- ↗ Weight decay (aka L2 regularization)
 - ↗ Penalizes the magnitude of the parameters
- ↗ Early stopping
 - ↗ Indirectly controls the magnitude of the parameters
- ↗ Alternatives
 - ↗ Dropout, batch normalization

Batch normalization



Batch normalization



Benefit: more stable and faster training. Often generalizes better

Against overfitting: data augmentation

- Resizing images keeping the aspect ratio.
- Enhancing images using random distortions (color, contrast, brightness and sharpness).
- Applying random crops using the same dimension for the width and height.
- Applying random horizontal flips.

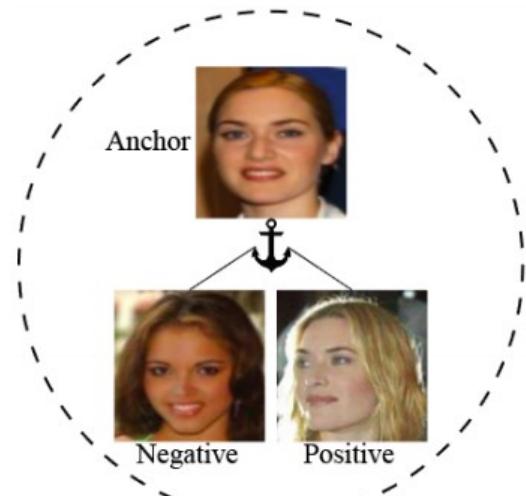


- ↗ AI, Machine learning & Deep learning
- ↗ What is a Convolutional Neural Network?
 - ↗ Layers
 - ↗ Optimization
- ↗ Applications

Convolutional NN

CNN beat humans in many tasks as:

- ↗ object recognition,
- ↗ lip reading,
- ↗ high-end surveillance,
- ↗ facial recognition,
- ↗ object-based searches,
- ↗ license plate readers,
- ↗ traffic violations detection,
- ↗ breast tomosynthesis diagnosis,
- ↗ etc., etc.



Biometric data for cheating detection

I

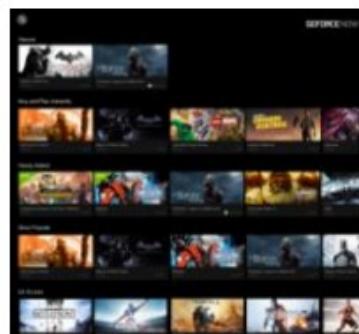
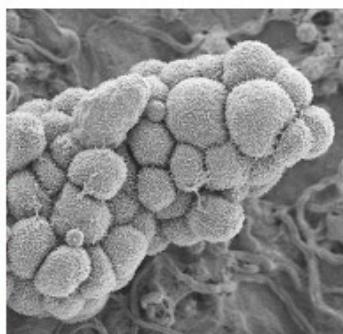


Analizes up to 38 faciales
facial micro-gestures, digital
prints and scans hand veins

iBorderCtrl will reduce the cost of
frontiers access of near 700M
crossing European countries



Deep learning everywhere



INTERNET & CLOUD

Image Classification
Speech Recognition
Language Translation
Language Processing
Sentiment Analysis
Recommendation

MEDICINE & BIOLOGY

Cancer Cell Detection
Diabetic Grading
Drug Discovery

MEDIA & ENTERTAINMENT

Video Captioning
Video Search
Real Time Translation

SECURITY & DEFENSE

Face Detection
Video Surveillance
Satellite Imagery

AUTONOMOUS MACHINES

Pedestrian Detection
Lane Tracking
Recognize Traffic Sign

Deep learning - one of the 10 breakthrough technologies

MIT Technology Review

10 BREAKTHROUGH TECHNOLOGIES 2013

Introduction The 10 Technologies Past Years

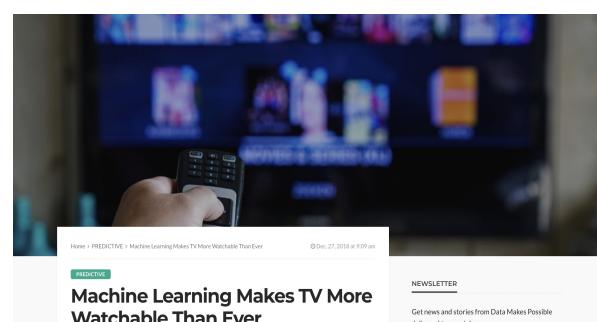
DeepLearning With massive amounts of computational power, machines can now recognize objects and translate speech in real time. Artificial intelligence is finally getting smart. →	Temporary Social Media Messages that quickly self-destruct could enhance the privacy of online communications and make people freer to be spontaneous. →	Prenatal DNA Sequencing Reading the DNA of fetuses will be the next frontier of the genomic revolution. But do you really want to know about the genetic problems or musical aptitude of your unborn child? →	Additive Manufacturing Skeptical about 3-D printing? GE, the world's largest manufacturer, is on the verge of using the technology to make jet parts. →	Baxter: The Blue-Collar Robot Rodney Brooks's newest creation is easy to interact with, but the complex innovations behind the robot show just how hard it is to get along with people. →
Memory Implants A maverick neuroscientist believes he has deciphered the code by which the brain forms long-term memories. Next: testing a prosthetic implant for people suffering from long-term memory loss. →	Smart Watches The designers of the Pebble watch realized that a mobile phone is more useful if you don't have to take it out of your pocket. →	Ultra-Efficient Solar Power Doubling the efficiency of a solar cell would completely change the economics of renewable energy. Nanotechnology just might make it possible. →	Big Data from Cheap Phones Collecting and analyzing information from simple cell phones can provide surprising insights into how people move about and behave – and even help us understand the spread of diseases. →	Supergrids A new high-power circuit breaker could finally make highly efficient DC power grids practical. →

"Preparing for the Future of Artificial Intelligence." ★
White House Frontiers Conference
THE WHITE HOUSE
FRONTIERS CONFERENCE
FRONTIERS CONFERENCE

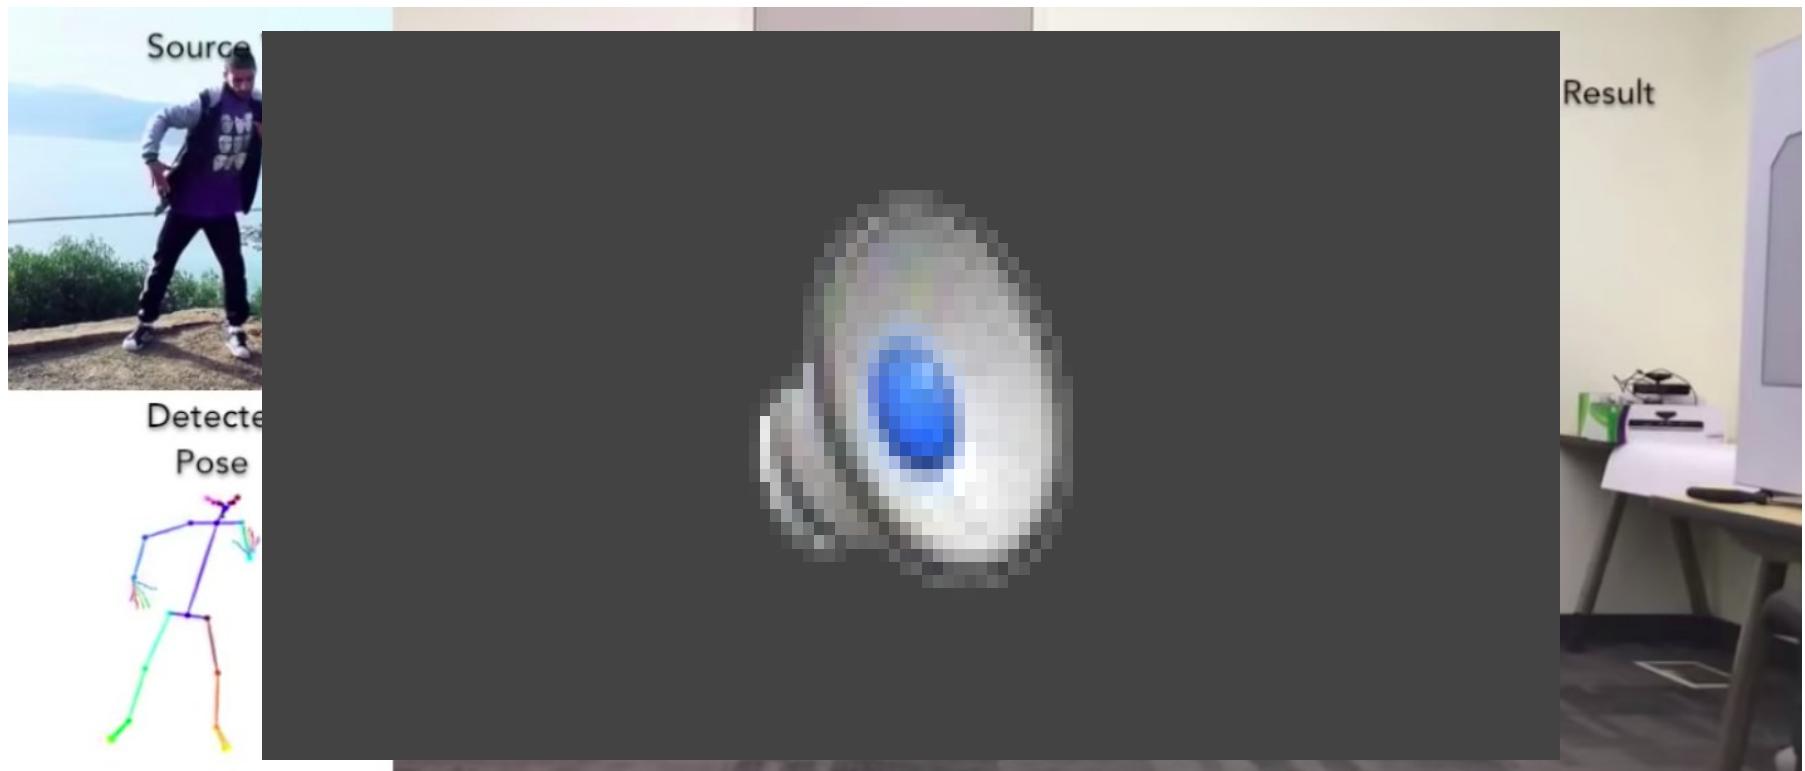


This 58-page report outlines a number of important topics related to artificial intelligence.

Deep learning everywhere



Everybody dances now



AI & ML: from R+D lab to real life



TOPICS ▾ MAIN MENU ☰

APAC EMEA Global Edition

FDA permits marketing of AI software that autonomously detects diabetic retinopathy

By [Dave Muio](#) | April 12, 2018 | 11:10 am

SHARE 733

The FDA has granted diagnostics company IDx's De Novo request to market its AI-based software system for the autonomous detection of diabetic retinopathy in adults who have diabetes, called IDx-DR.



This decision represents the first AI-based diagnostic system authorized by the FDA for commercialization in the US that can provide a screening decision without the need for clinician interpretation, according to the agency. The news comes just months after the [FDA's De Novo approval of Viz.AI](#), another AI software tool that analyzes stroke indicators and highlights CT images that could require additional clinical attention.

"Early detection of retinopathy is an important part of managing care for the millions of people with diabetes, yet many patients with diabetes are not adequately screened for diabetic retinopathy since about 50 percent of them do not see their eye doctor on a yearly basis," Dr. Malvina Eydelman, director of the Division of Ophthalmic and Ear, Nose, and Throat Devices at the FDA's Center for Devices and Radiological Health, said in a statement. "Today's decision permits the marketing of a novel artificial



(China Daily) 14:15, July 02, 2018

Follow on Apple News



Radiologist Zhang Junhai from Shanghai Huashan Hospital reads a medical image display during a competition with BioMind, an artificial intelligence system, in Beijing on Saturday. (Photo by Chen Zihao/China Daily)

CNN MARKETS BUSINESS INVESTING TECH POLITICS CNBC-TV WATCHLIST PRO ▾

Google's DeepMind A.I. beats doctors in breast cancer screening trial

PUBLISHED THU, JAN 2 2020 8:13 AM EST | UPDATED THU, JAN 2 2020 8:13 AM EST

David Reid [@DAVIDREID73](#)

SHARE

- KEY POINTS
- Anonymous scans of 29,000 women were used in the trial.
 - The biggest improvements over human scanning was found in the U.S. portion of the study.
 - Google-owned DeepMind has already used AI to read eye scans and spot neck cancer.



Become part of the Revolut

Get trained, get connected and... by the most sou graduate emplo

AmplifyME

Learn more >

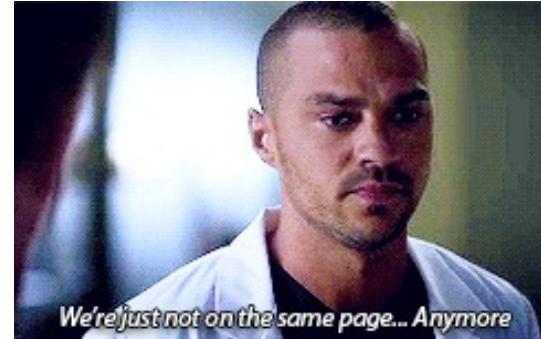
TRENDING NOW

01:42

•AI & ML: from R+D lab to hospital

Concerns about DL

- Different development approach.



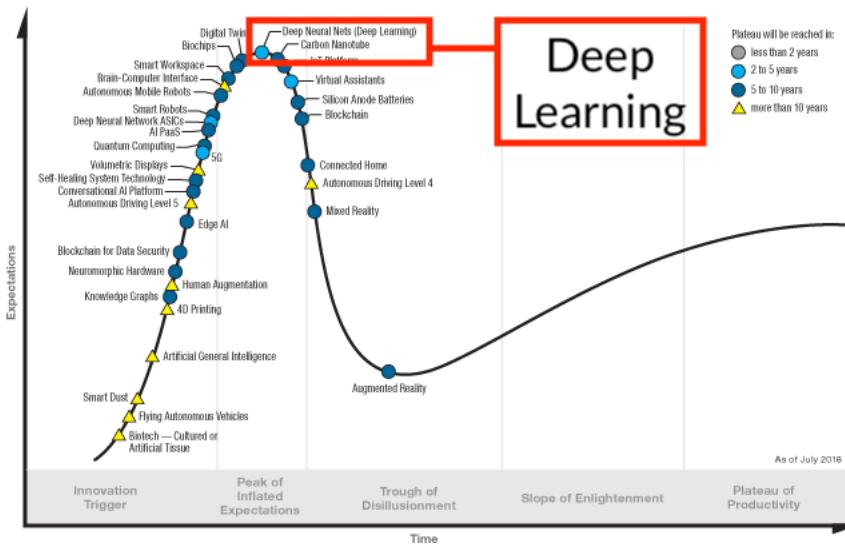
- No clear view on how insight is generated.



- A system is only as good as the data it learns from.

DL and society expectation

Hype Cycle for Emerging Technologies, 2018



gartner.com/SmarterWithGartner

Source: Gartner (August 2018)
© 2018 Gartner, Inc. and/or its affiliates. All rights reserved.

Gartner

Hype Cycle for Artificial Intelligence, 2022



gartner.com

Source: Gartner
© 2022 Gartner, Inc. and/or its affiliates. All rights reserved. Gartner and Hype Cycle are registered trademarks of Gartner, Inc. and its affiliates in the U.S. 1957302

Gartner

Deep Learning's 'Permanent Peak' On Gartner's Hype Cycle

Conclusions

|

Deep learning – a technology that came to stay

A (new) technological trend that is affecting directly our environment
Multiple applications

Different CNN models can be found. No optimal one exists.

The most important elements:

- Layers
- Loss function
- Fine-tunning, optimization
- Overfitting

CNN applicable to most Computer Vision problems

01:42

Trends in DL



1. Transfer Learning: The use of pre-trained models and transfer learning techniques was becoming more widespread, as they can significantly reduce the amount of data required to train effective models.

- Multi-task learning

2. Uncertainty modeling: refers to the process of quantifying and managing uncertainty or ambiguity in the predictions or decisions made by machine learning models

- Uncertainty-aware MTL
- Learning with noisy labeling

3. Self-Supervised Learning: methods train models on unlabeled data, which can be particularly useful in cases where labeled data is scarce.

4. Meta-Learning: Meta-learning techniques were explored to enable models to learn how to learn, which can lead to faster adaptation to new tasks.

- Continual learning

5. Generative AI and Generative Adversarial Networks (GANs): GANs were being used for a variety of applications, from image generation to data augmentation and domain adaptation.

- Uncertainty-aware data augmentation
- NeRFs
- Stable diffusion

Trends in DL

6. Efficiency and Model Compression: There was a growing interest in making deep learning models smaller, faster, and more energy-efficient, particularly for edge computing and mobile applications.

- Scaling by hierarchical DL models
- Fine-grained recognition

7. Multimodal Learning: Combining information from different sources, such as text and images, to create more comprehensive models for understanding and generating content.

- Oncology, Food ontology

8. Explainable AI (XAI): The need for understanding and interpreting deep learning models became more critical, especially in fields like healthcare and finance where model explainability is crucial.

- Robust Explainable models

9. Federated Learning: This approach allows for training models across decentralized data sources without sharing raw data, preserving privacy. It gained traction, particularly in applications involving sensitive or proprietary data.

10. AI for Healthcare: Deep learning was increasingly applied to medical image analysis, drug discovery, and patient diagnosis, with a focus on improving healthcare outcomes.

- Our projects: food data analysis, omics data,

